

ml-project

October 26, 2025

```
[26]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from numpy import array
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
```

```
[27]: from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

```
[28]: !ls
```

drive sample_data sample_submission.csv test.csv train.csv

```
[29]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[30]: trainData=pd.read_csv("train.csv")
trainData
```

```
[30]:
```

	id	Gender	Age	Height	Weight	\
0	0	Male	24.443011	1.699998	81.669950	
1	1	Female	18.000000	1.560000	57.000000	
2	2	Female	18.000000	1.711460	50.165754	

3	3	Female	20.952737	1.710730	131.274851
4	4	Male	31.641081	1.914186	93.798055
...
15528	15528	Male	18.000000	1.700000	50.000000
15529	15529	Male	18.000000	1.763101	55.523481
15530	15530	Female	19.010211	1.686936	49.660995
15531	15531	Male	22.777890	1.805445	85.228116
15532	15532	Male	39.371523	1.770278	79.677930

	family_history_with_overweight	FAVC	FCVC	NCP	CAEC \
0	yes	yes	2.000000	2.983297	Sometimes
1	yes	yes	2.000000	3.000000	Frequently
2	yes	yes	1.880534	1.411685	Sometimes
3	yes	yes	3.000000	3.000000	Sometimes
4	yes	yes	2.679664	1.971472	Sometimes
...
15528	no	yes	2.000000	3.000000	Frequently
15529	yes	yes	2.786008	3.000000	Sometimes
15530	no	yes	1.053534	3.452590	Sometimes
15531	yes	yes	2.000000	2.092179	Sometimes
15532	yes	yes	2.407817	1.097312	Sometimes

	SMOKE	CH2O	SCC	FAF	TUE	CALC \
0	no	2.763573	no	0.000000	0.976473	Sometimes
1	no	2.000000	no	1.000000	1.000000	no
2	no	1.910378	no	0.866045	1.673584	no
3	no	1.674061	no	1.467863	0.780199	Sometimes
4	no	1.979848	no	1.967973	0.931721	Sometimes
...
15528	no	2.000000	no	1.000000	2.000000	Sometimes
15529	no	1.962646	yes	0.028202	1.561272	Sometimes
15530	no	1.000000	no	2.001230	1.000000	Sometimes
15531	no	2.452986	no	0.796770	0.000000	Sometimes
15532	no	2.205911	no	0.977929	0.000000	Frequently

	MTRANS	WeightCategory
0	Public_Transportation	Overweight_Level_II
1	Automobile	Normal_Weight
2	Public_Transportation	Insufficient_Weight
3	Public_Transportation	Obesity_Type_III
4	Public_Transportation	Overweight_Level_II
...
15528	Public_Transportation	Insufficient_Weight
15529	Public_Transportation	Insufficient_Weight
15530	Public_Transportation	Insufficient_Weight
15531	Public_Transportation	Overweight_Level_I
15532	Automobile	Overweight_Level_II

[15533 rows x 18 columns]

```
[31]: trainData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15533 entries, 0 to 15532
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    15533 non-null  int64
1   Gender                               15533 non-null  object
2   Age                                   15533 non-null  float64
3   Height                               15533 non-null  float64
4   Weight                               15533 non-null  float64
5   family_history_with_overweight       15533 non-null  object
6   FAVC                                  15533 non-null  object
7   FCVC                                  15533 non-null  float64
8   NCP                                   15533 non-null  float64
9   CAEC                                  15533 non-null  object
10  SMOKE                                 15533 non-null  object
11  CH2O                                  15533 non-null  float64
12  SCC                                   15533 non-null  object
13  FAF                                   15533 non-null  float64
14  TUE                                   15533 non-null  float64
15  CALC                                  15533 non-null  object
16  MTRANS                                15533 non-null  object
17  WeightCategory                       15533 non-null  object
dtypes: float64(8), int64(1), object(9)
memory usage: 2.1+ MB
```

```
[32]: trainData.columns
```

```
[32]: Index(['id', 'Gender', 'Age', 'Height', 'Weight',
          'family_history_with_overweight', 'FAVC', 'FCVC', 'NCP', 'CAEC',
          'SMOKE', 'CH2O', 'SCC', 'FAF', 'TUE', 'CALC', 'MTRANS',
          'WeightCategory'],
          dtype='object')
```

```
[33]: trainData.dtypes
```

```
[33]: id                                int64
      Gender                           object
      Age                               float64
      Height                           float64
      Weight                           float64
      family_history_with_overweight    object
```

```

FAVC          object
FCVC          float64
NCP           float64
CAEC          object
SMOKE         object
CH2O          float64
SCC           object
FAF           float64
TUE           float64
CALC          object
MTRANS        object
WeightCategory object
dtype: object

```

```
[34]: trainData.describe()
```

```

[34]:
count    15533.000000    15533.000000    15533.000000    15533.000000    15533.000000 \
mean      7766.000000      23.816308      1.699918      87.785225      2.442917
std       4484.135201       5.663167      0.087670      26.369144      0.530895
min         0.000000      14.000000      1.450000      39.000000      1.000000
25%       3883.000000      20.000000      1.630927      66.000000      2.000000
50%       7766.000000      22.771612      1.700000      84.000000      2.342220
75%      11649.000000      26.000000      1.762921     111.600553      3.000000
max      15532.000000      61.000000      1.975663     165.057269      3.000000

count    15533.000000    15533.000000    15533.000000    15533.000000
mean         2.760425      2.027626      0.976968      0.613813
std         0.706463      0.607733      0.836841      0.602223
min         1.000000      1.000000      0.000000      0.000000
25%         3.000000      1.796257      0.007050      0.000000
50%         3.000000      2.000000      1.000000      0.566353
75%         3.000000      2.531456      1.582675      1.000000
max         4.000000      3.000000      3.000000      2.000000

```

```
[35]: trainData.describe(include='object')
```

```

[35]:
count    15533    15533    15533    15533    15533    15533 \
unique         2         2         2         4         2         2
top      Male     yes     yes  Sometimes     no     no
freq      7783    12696    14184     13126    15356    15019

count    15533    15533    15533
unique         3         5         7

```

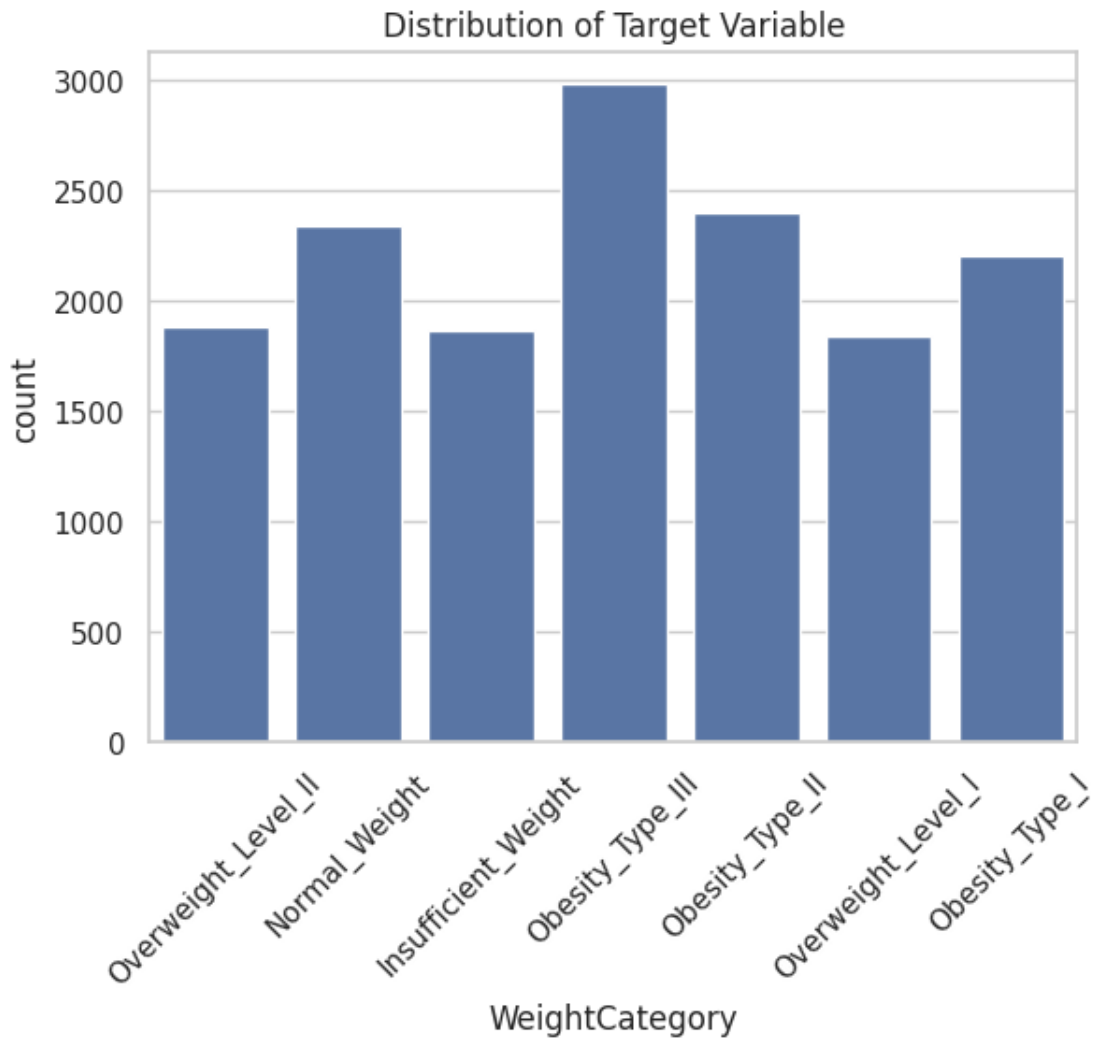
top	Sometimes	Public_Transportation	Obesity_Type_III
freq	11285	12470	2983

1 EXPLORATORY DATA ANALYSIS

Target Variable Analysis

```
[36]: sns.countplot(x='WeightCategory', data=trainData)
plt.title("Distribution of Target Variable")
plt.xticks(rotation=45)
plt.show()

trainData['WeightCategory'].value_counts(normalize=True)
```

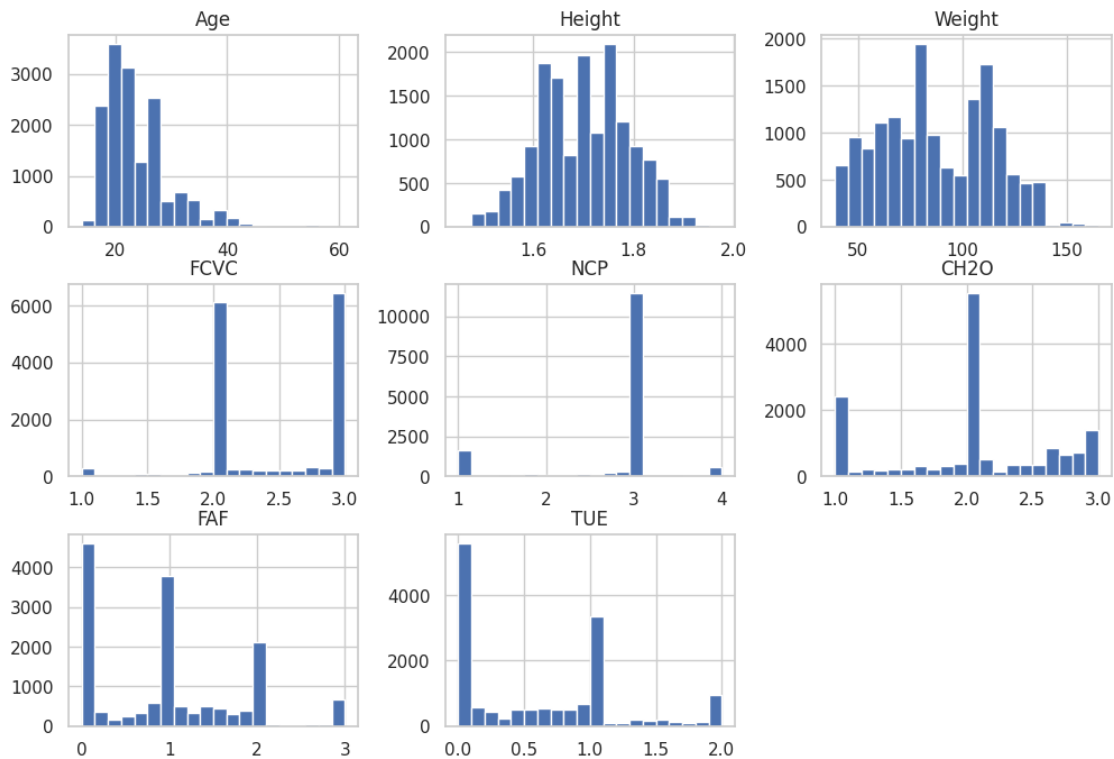


```
[36]: WeightCategory
Obesity_Type_III      0.192043
Obesity_Type_II       0.154703
Normal_Weight         0.150969
Obesity_Type_I        0.142085
Overweight_Level_II   0.121097
Insufficient_Weight   0.120389
Overweight_Level_I    0.118715
Name: proportion, dtype: float64
```

The target variable WeightCategory shows a fairly uniform distribution across all seven categories, with class proportions ranging from 11.9% to 19.2%. The maximum-to-minimum ratio is approximately 1.6:1, indicating that the dataset is reasonably balanced. Hence, no major imbalance handling is required.

Univariate Analysis for Numerical Variables

```
[37]: numeric_cols = ['Age', 'Height', 'Weight', 'FCVC', 'NCP', 'CH2O', 'FAF', 'TUE']
trainData[numeric_cols].hist(bins=20, figsize=(12, 8))
plt.show()
```



- The **Age** distribution is **right-skewed**, with the majority of participants falling between **18 and 25 years**, indicating that the dataset is dominated by young adults.

- **Height** follows an approximately **normal distribution**, centered around **1.65 to 1.75 meters**, suggesting a realistic spread across participants.
- **Weight** is **moderately right-skewed**, with most individuals weighing between **50 and 100 kilograms**, and a few higher values that likely represent genuinely obese participants.
- **FCVC (Frequency of Vegetable Consumption)** shows a **bimodal pattern**, where most individuals have values around **2 or 3**, implying regular vegetable consumption among participants.
- **NCP (Number of Main Meals)** displays a strong concentration at **3 meals per day**, which aligns with typical eating habits.
- **CH2O (Daily Water Intake)** is **mildly right-skewed**, with most participants consuming about **2 liters of water per day**.
- **FAF (Physical Activity Frequency)** has a **multimodal distribution**, suggesting a wide range of exercise habits, from no activity to frequent exercise.
- **TUE (Time Using Technology Devices)** is **highly right-skewed**, showing that most participants spend **less than one hour per day** using technology.

A few outliers are present in Age and Weight, but they appear to be genuine rather than data entry errors.

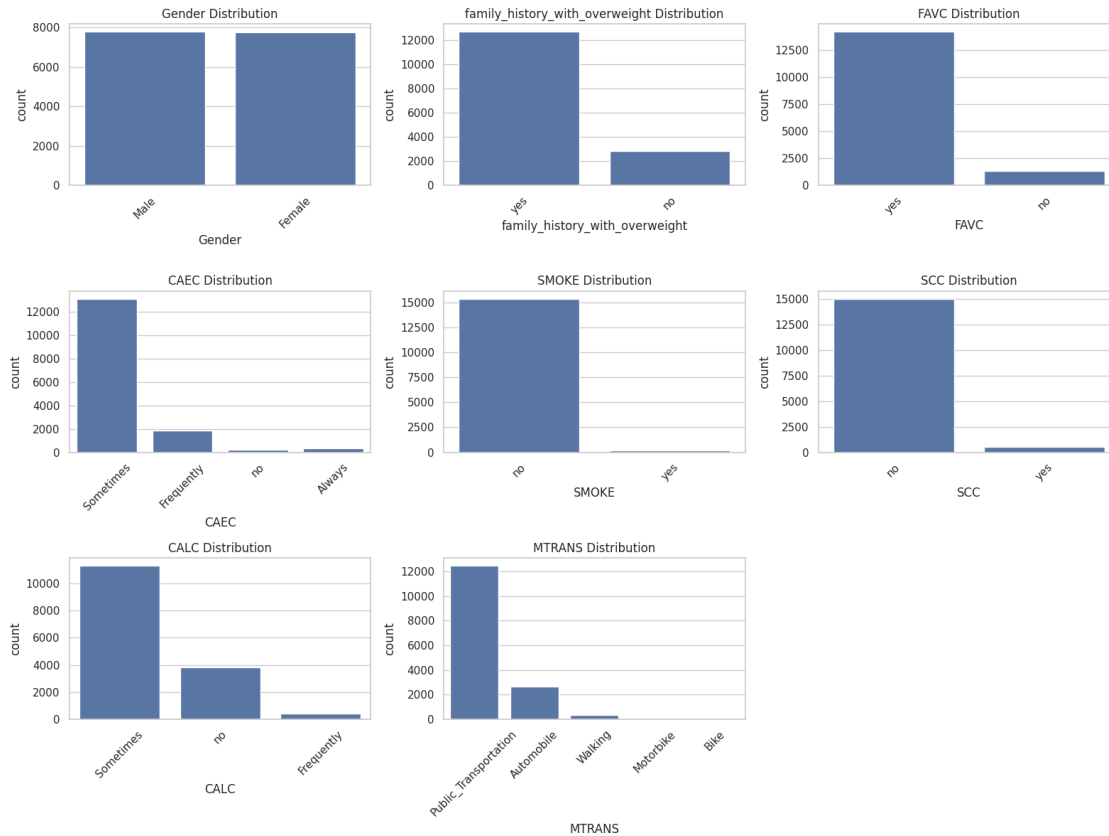
Overall, all numerical features show logical and interpretable distributions, with only mild skewness in some variables that can be handled during model preparation.

Univariate Analysis for Categorical Variables

```
[38]: cat_cols = ['Gender', 'family_history_with_overweight', 'FAVC', 'CAEC',
                  'SMOKE', 'SCC', 'CALC', 'MTRANS']

plt.figure(figsize=(16, 12))
for i, col in enumerate(cat_cols, 1):
    plt.subplot(3, 3, i)
    sns.countplot(x=col, data=trainData)
    plt.title(f"{col} Distribution")
    plt.xticks(rotation=45)
    plt.tight_layout()

plt.show()
```



```
[39]: cat_cols = ['Gender', 'family_history_with_overweight', 'FAVC', 'CAEC',
                  'SMOKE', 'SCC', 'CALC', 'MTRANS']
```

```
for col in cat_cols:
    print(f"\n{'='*40}")
    print(f" {col} Distribution")
    print(f"{'='*40}")
    counts = trainData[col].value_counts(normalize=True) * 100
    for cat, val in counts.items():
        print(f"{cat:<25}: {val:5.2f}%")
```

```
=====
Gender Distribution
=====
Male                : 50.11%
Female              : 49.89%

=====
family_history_with_overweight Distribution
=====
```


yes	: 81.74%
no	: 18.26%

=====

FAVC Distribution

=====

yes	: 91.32%
no	: 8.68%

=====

CAEC Distribution

=====

Sometimes	: 84.50%
Frequently	: 11.96%
Always	: 2.23%
no	: 1.31%

=====

SMOKE Distribution

=====

no	: 98.86%
yes	: 1.14%

=====

SCC Distribution

=====

no	: 96.69%
yes	: 3.31%

=====

CALC Distribution

=====

Sometimes	: 72.65%
no	: 24.73%
Frequently	: 2.62%

=====

MTRANS Distribution

=====

Public_Transportation	: 80.28%
Automobile	: 17.18%
Walking	: 2.19%
Motorbike	: 0.19%
Bike	: 0.15%

- The **Gender distribution is nearly equal**, with 50.11% males and 49.89% females, indicating a perfectly balanced gender representation in the dataset.
- A large proportion (**81.74%**) of individuals reported having a family history of over-

weight, suggesting a strong hereditary influence among participants.

- The **majority (91.32%) of participants responded “yes” to FAVC (Frequent Consumption of High-Calorie Food)**, implying that most individuals regularly consume calorie-dense foods.
- For **CAEC (Consumption of Food Between Meals)**, **84.5% of respondents reported “Sometimes”**, while only 1.31% reported “no,” indicating that snacking between meals is a common habit.
- **SMOKE shows a clear pattern**, with **98.86% non-smokers** and only 1.14% smokers, suggesting that smoking is relatively uncommon in this population.
- In **SCC (Calories Monitoring)**, **96.69% of individuals do not monitor their calorie intake**, showing that calorie tracking is rare among participants.
- For **CALC (Alcohol Consumption)**, **72.65% reported drinking “Sometimes”**, while 24.73% said “no”, indicating that moderate alcohol consumption is fairly common.
- Regarding **MTRANS (Mode of Transportation)**, the **vast majority (80.28%) use public transportation**, followed by automobiles (17.18%), while walking and biking are much less common.

Overall Insights

Most categorical variables are highly imbalanced, with one dominant category (e.g., FAVC = “yes”, SMOKE = “no”).

However, these distributions reflect realistic lifestyle behaviors rather than data issues.

Variables such as Gender are well-balanced, while others (e.g., SMOKE, SCC, MTRANS) may provide limited variability for model training.

The categorical data appears clean, with no evidence of inconsistent or erroneous category values.

Bivariate Analysis for Numerical Variables

```
[40]: sns.set_theme(style="whitegrid")

plt.figure(figsize=(16, 12))

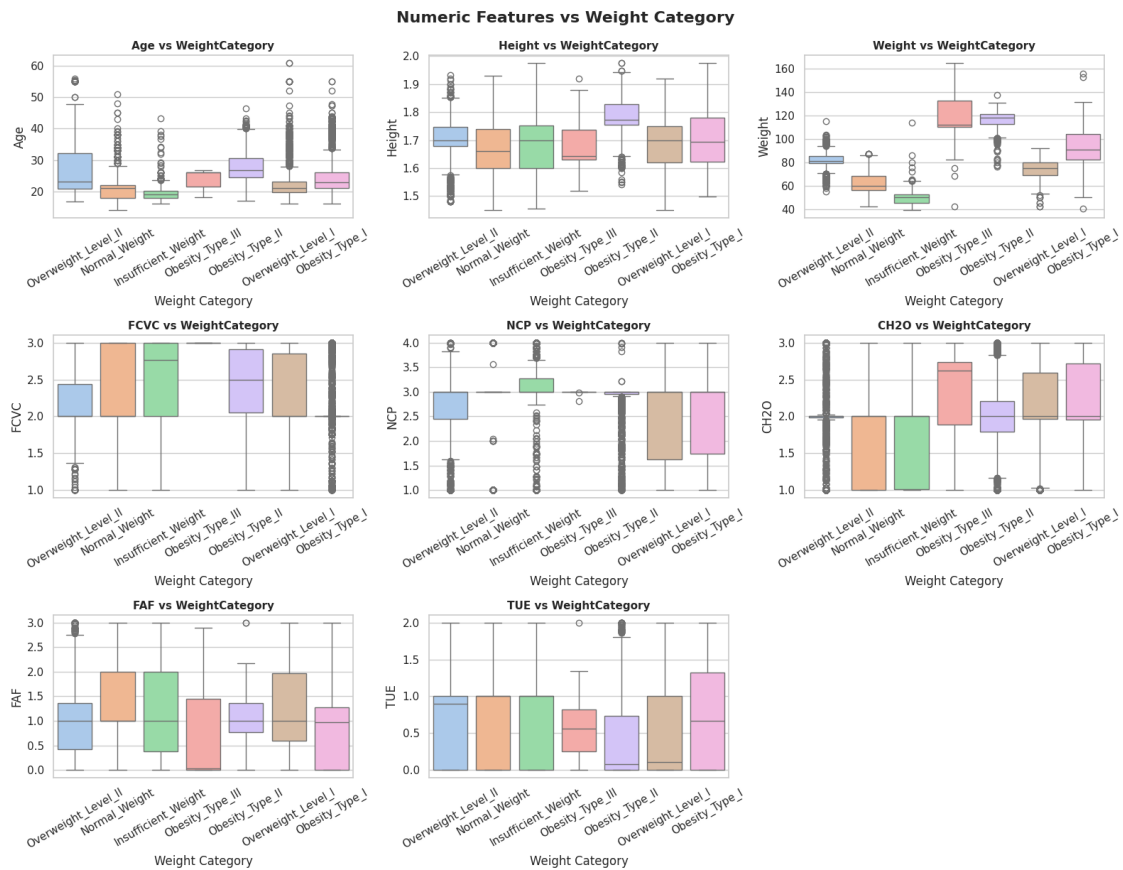
for i, col in enumerate(numeric_cols, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(
        x='WeightCategory',
        y=col,
        hue='WeightCategory',
        data=trainData,
        palette='pastel',
        legend=False
    )
    plt.title(f"{col} vs WeightCategory", fontsize=11, weight='bold')
    plt.xlabel("Weight Category")
```

```

plt.ylabel(col)
plt.xticks(rotation=30)
plt.tight_layout()

plt.suptitle("Numeric Features vs Weight Category", fontsize=16, weight='bold',
             y=1.02)
plt.show()

```



Bivariate Analysis for Categorical Variables

```

[41]: import math

sns.set_theme(style="whitegrid")

n_cols = 3
n_rows = math.ceil(len(cat_cols) / n_cols)

plt.figure(figsize=(6 * n_cols, 4 * n_rows))

for i, col in enumerate(cat_cols, 1):

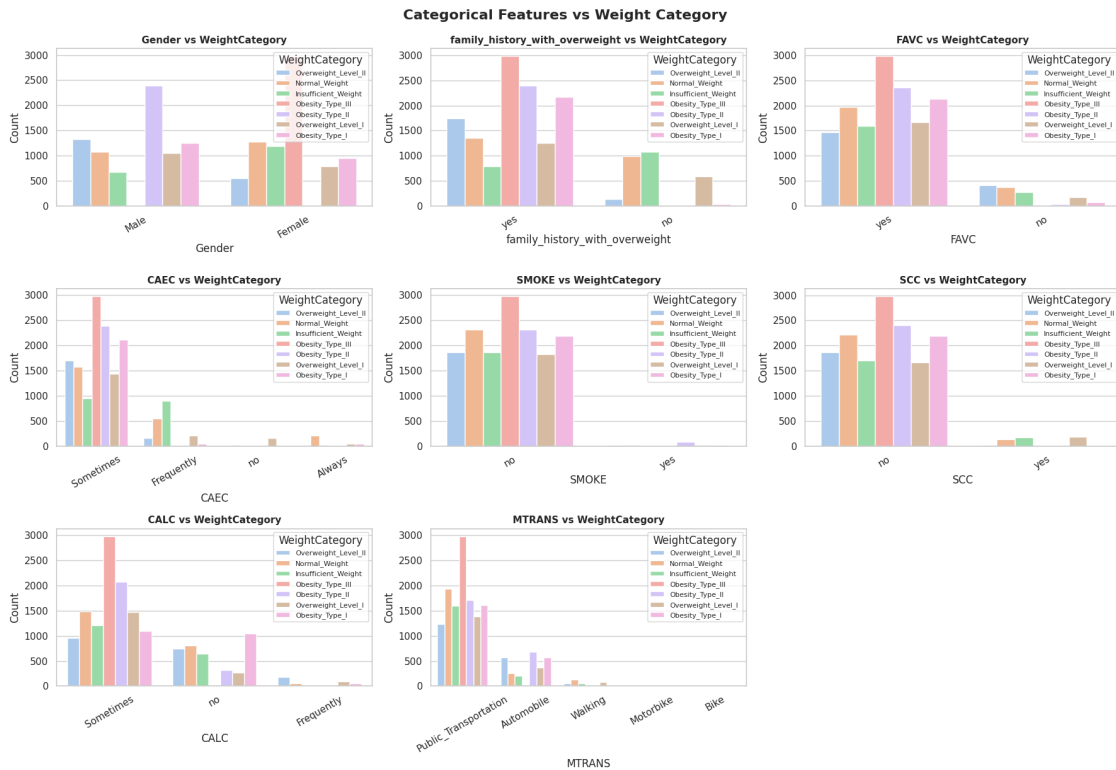
```

```

plt.subplot(n_rows, n_cols, i)
sns.countplot(
    x=col,
    hue='WeightCategory',
    data=trainData,
    palette='pastel'
)
plt.title(f"{col} vs WeightCategory", fontsize=11, weight='bold')
plt.xlabel(col)
plt.ylabel("Count")
plt.xticks(rotation=30)
plt.legend(title="WeightCategory", loc="upper right", fontsize=8)
plt.tight_layout()

plt.suptitle("Categorical Features vs Weight Category", fontsize=16,
    weight='bold', y=1.02)
plt.show()

```



Correlation Analysis

```

[42]: x=trainData.drop('WeightCategory',axis=1)
target=trainData["WeightCategory"]
target.head()

```


```

df_copy = x.copy()
categorical_cols = df_copy.select_dtypes(include='object').columns

for col in categorical_cols:
    df_copy[col] = LabelEncoder().fit_transform(df_copy[col])

plt.figure(figsize=(15,10))
sns.heatmap(df_copy.corr(), annot=True, cmap='coolwarm')
plt.show()

```



The above heatmap depicts the correlation between the features. There are no two features having high correlation between them. Hence none of the features were dropped.


Outlier Detection

```

[43]: plt.figure(figsize=(12,8))
for i,col in enumerate(numeric_cols,1):
    plt.subplot(2,4,i)
    sns.boxplot(trainData[col])
    plt.title(col)
plt.tight_layout()

```

```
plt.show()
```



```
[44]: # IQR method to flag outliers
def outlier_indices(series):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5*IQR
    upper = Q3 + 1.5*IQR
    return series[(series < lower) | (series > upper)].index

outlier_dict = {col: outlier_indices(trainData[col]) for col in numeric_cols}

# check which columns have too many outliers
for col, idx in outlier_dict.items():
    print(f"{col}: {len(idx)} outliers")
```

```
Age: 792 outliers
Height: 4 outliers
Weight: 0 outliers
FCVC: 0 outliers
NCP: 4548 outliers
CH2O: 0 outliers
FAF: 0 outliers
```

TUE: 0 outliers

```
[45]: # Age had 792 outliers but its Maximum value was 61 and Minimum value was 14
      # which can't be ignored.
```

```
# Compute IQR for Height
```

```
Q1 = trainData['Height'].quantile(0.25)
```

```
Q3 = trainData['Height'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
lower = Q1 - 1.5*IQR
```

```
upper = Q3 + 1.5*IQR
```

```
# Show the outlier values
```

```
trainData['Height'][(trainData['Height'] < lower) | (trainData['Height'] >
    ↪upper)]
```

```
#Conclulsion - The outlier heights are under 2m. Hence these can't be called
# bad outliers.
```

```
[45]: 1271      1.975663
      4084      1.975663
      7272      1.975663
      11718     1.975663
      Name: Height, dtype: float64
```

2 Data Preprocessing

```
[46]: trainData.isnull().any()
```

```
[46]: id                False
      Gender            False
      Age               False
      Height            False
      Weight            False
      family_history_with_overweight  False
      FAVC              False
      FCVC              False
      NCP               False
      CAEC              False
      SMOKE             False
      CH20              False
      SCC               False
      FAF               False
      TUE               False
      CALC              False
```

```
MTRANS                                False
WeightCategory                        False
dtype: bool
```

```
[47]: trainData.duplicated().any()
```

```
[47]: np.False_
```

```
[48]: x=trainData.drop('WeightCategory',axis=1)
      target=trainData["WeightCategory"]
      target.head()
```

```
[48]: 0    Overweight_Level_II
      1      Normal_Weight
      2    Insufficient_Weight
      3      Obesity_Type_III
      4    Overweight_Level_II
      Name: WeightCategory, dtype: object
```

```
[49]: x.head()
```

```
[49]:   id  Gender      Age      Height      Weight family_history_with_overweight \
0   0   Male  24.443011  1.699998   81.669950                      yes
1   1  Female  18.000000  1.560000   57.000000                      yes
2   2  Female  18.000000  1.711460   50.165754                      yes
3   3  Female  20.952737  1.710730  131.274851                      yes
4   4   Male  31.641081  1.914186   93.798055                      yes
```

```
      FAVC      FCVC      NCP      CAEC SMOKE      CH20 SCC      FAF \
0  yes  2.000000  2.983297  Sometimes    no  2.763573  no  0.000000
1  yes  2.000000  3.000000  Frequently    no  2.000000  no  1.000000
2  yes  1.880534  1.411685  Sometimes    no  1.910378  no  0.866045
3  yes  3.000000  3.000000  Sometimes    no  1.674061  no  1.467863
4  yes  2.679664  1.971472  Sometimes    no  1.979848  no  1.967973
```

```
      TUE      CALC      MTRANS
0  0.976473  Sometimes  Public_Transportation
1  1.000000         no      Automobile
2  1.673584         no  Public_Transportation
3  0.780199  Sometimes  Public_Transportation
4  0.931721  Sometimes  Public_Transportation
```

```
[50]: numeric_cols = x.select_dtypes(include=[np.number]).columns.tolist()
      categorical_cols = x.select_dtypes(include=['object', 'category']).columns.
      ↪tolist()
      from sklearn.preprocessing import OneHotEncoder
      X_encoded = pd.get_dummies(x, drop_first=True) # simple one-hot encoding
```


LOADING DATASET FOR TRAINING AND TESTING MODELS

```
[51]: trainData = pd.read_csv("train.csv")
      testData = pd.read_csv("test.csv")

      X = trainData.drop(columns=['WeightCategory', 'id'])
      y = trainData['WeightCategory']

      X_test = testData.drop(columns=['id'])
```

1. DECISION TREE

```
[52]: cat_cols = [
      ↪['Gender', 'family_history_with_overweight', 'FAVC', 'CAEC', 'SMOKE', 'SCC', 'CALC', 'MTRANS']]
      le = LabelEncoder()
      for col in cat_cols:
          x[col] = le.fit_transform(x[col])
      X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42)
```

```
[53]: dt_clf = DecisionTreeClassifier(random_state=42)

      dt_clf.fit(X_train, y_train)

      y_pred = dt_clf.predict(X_test)

      dt_clf = DecisionTreeClassifier(
          criterion='gini',
          max_depth=7,
          min_samples_split=5,
          min_samples_leaf=2,
          random_state=42
      )

      dt_clf.fit(X_train, y_train)

      y_pred = dt_clf.predict(X_test)

      test_accuracy_dtt = accuracy_score(y_test, y_pred)
      test_precision_dtt = precision_score(y_test, y_pred, average='weighted',
      ↪zero_division=0)
      test_recall_dtt = recall_score(y_test, y_pred, average='weighted',
      ↪zero_division=0)
      test_f1_dtt = f1_score(y_test, y_pred, average='weighted', zero_division=0)

      print("--- Decision Tree Model Evaluation ---")
      print("Test Accuracy:", test_accuracy_dtt)
      print("Test Precision (Weighted):", test_precision_dtt)
```

```

print("Test Recall (Weighted):", test_recall_dtt)
print("Test F1-Score (Weighted):", test_f1_dtt)
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

importances = dt_clf.feature_importances_

feature_names = X_train.columns

feat_importances = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'],
        color='skyblue')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from Decision Tree")
plt.show()

```

--- Decision Tree Model Evaluation ---

Test Accuracy: 0.8651432249758609
 Test Precision (Weighted): 0.8659294193764477
 Test Recall (Weighted): 0.8651432249758609
 Test F1-Score (Weighted): 0.8654772440534064

Confusion Matrix:

```


[[307  36   1   1   0   0   1]
 [ 29 413   1   0   0  50   3]
 [  0   0 393  15   1  10  29]
 [  0   0  26 436   0   0   2]
 [  0   0   3   1 602   0   0]
 [  3  49   8   0   0 268  56]
 [  0   6  27   1   0  60 269]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.91	0.89	0.90	346
Normal_Weight	0.82	0.83	0.83	496
Obesity_Type_I	0.86	0.88	0.87	448
Obesity_Type_II	0.96	0.94	0.95	464
Obesity_Type_III	1.00	0.99	1.00	606

Overweight_Level_I	0.69	0.70	0.69	384
Overweight_Level_II	0.75	0.74	0.74	363
accuracy			0.87	3107
macro avg	0.85	0.85	0.85	3107
weighted avg	0.87	0.87	0.87	3107



ACCURACY WITH DECISION TREE IS 0.8651432249758609

Hyperparameter Tuning Using Optuna on Decision Tree

```
[54]: !pip install optuna
```

```
Collecting optuna
  Downloading optuna-4.5.0-py3-none-any.whl.metadata (17 kB)
Collecting alembic>=1.5.0 (from optuna)
  Downloading alembic-1.17.0-py3-none-any.whl.metadata (7.2 kB)
Collecting colorlog (from optuna)
  Downloading colorlog-6.10.1-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from optuna) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from optuna) (25.0)
Collecting sqlalchemy>=1.4.2 (from optuna)
  Downloading sqlalchemy-2.0.44-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (9.5 kB)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from optuna) (4.67.1)
```

Requirement already satisfied: PyYAML in /usr/local/lib/python3.12/dist-packages (from optuna) (6.0.3)

Requirement already satisfied: Mako in /usr/lib/python3/dist-packages (from alembic>=1.5.0->optuna) (1.1.3)

Requirement already satisfied: typing-extensions>=4.12 in /usr/local/lib/python3.12/dist-packages (from alembic>=1.5.0->optuna) (4.15.0)

Collecting greenlet>=1 (from sqlalchemy>=1.4.2->optuna)

Downloading greenlet-3.2.4-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl.metadata (4.1 kB)

Downloading optuna-4.5.0-py3-none-any.whl (400 kB)

400.9/400.9 kB

12.4 MB/s eta 0:00:00

Downloading alembic-1.17.0-py3-none-any.whl (247 kB)

247.4/247.4 kB

24.5 MB/s eta 0:00:00

Downloading sqlalchemy-2.0.44-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.3 MB)

3.3/3.3 MB

100.3 MB/s eta 0:00:00

Downloading colorlog-6.10.1-py3-none-any.whl (11 kB)

Downloading greenlet-3.2.4-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl (607 kB)

607.6/607.6 kB

47.9 MB/s eta 0:00:00

Installing collected packages: greenlet, colorlog, sqlalchemy, alembic, optuna

Successfully installed alembic-1.17.0 colorlog-6.10.1 greenlet-3.2.4 optuna-4.5.0 sqlalchemy-2.0.44

```
[55]: import optuna
```

```
[56]: def objective(trial):
    criterion = trial.suggest_categorical('criterion', ['gini', 'entropy', 'log_loss'])
    max_depth = trial.suggest_int('max_depth', 2, 20)
    min_samples_split = trial.suggest_int('min_samples_split', 2, 20)
    min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 10)
    dt_clf = DecisionTreeClassifier(
        criterion=criterion,
        max_depth=max_depth,
        min_samples_split=min_samples_split,
        min_samples_leaf=min_samples_leaf,
        random_state=42
    )
```

```

    score = cross_val_score(dt_clf, X_train, y_train, cv=5, scoring='accuracy').
    ↪mean()
    return score

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50, show_progress_bar=True)

best_params = study.best_params
dt_clf_best = DecisionTreeClassifier(**best_params, random_state=42)
dt_clf_best.fit(X_train, y_train)
y_pred = dt_clf_best.predict(X_test)

test_accuracy_dtt_hp = accuracy_score(y_test, y_pred)
test_precision_dtt_hp = precision_score(y_test, y_pred, average='weighted', ↪
    ↪zero_division=0)
test_recall_dtt_hp = recall_score(y_test, y_pred, average='weighted', ↪
    ↪zero_division=0)
test_f1_dtt_hp = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("Best parameters:", study.best_params)
print("Best cross-validation accuracy:", study.best_value)
print("--- Decision Tree Model Evaluation (Optimized) ---")
print("Test Accuracy:", test_accuracy_dtt_hp)
print("Test Precision (Weighted):", test_precision_dtt_hp)
print("Test Recall (Weighted):", test_recall_dtt_hp)
print("Test F1-Score (Weighted):", test_f1_dtt_hp)
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

importances = dt_clf_best.feature_importances_
feature_names = X_train.columns
feat_importances = pd.DataFrame({'Feature': feature_names, 'Importance': ↪
    ↪importances}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'], ↪
    ↪color='skyblue')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from Optimized Decision Tree")
plt.show()

```

[I 2025-10-26 16:26:14,114] A new study created in memory with name: no-name-e4d43aab-97e3-42da-bdb1-b5154228b207

0% | 0/50 [00:00<?, ?it/s]

[I 2025-10-26 16:26:14,530] Trial 0 finished with value: 0.8670532608361352 and

parameters: {'criterion': 'log_loss', 'max_depth': 16, 'min_samples_split': 13, 'min_samples_leaf': 9}. Best is trial 0 with value: 0.8670532608361352.

[I 2025-10-26 16:26:14,901] Trial 1 finished with value: 0.8665707195708443 and parameters: {'criterion': 'gini', 'max_depth': 15, 'min_samples_split': 20, 'min_samples_leaf': 6}. Best is trial 0 with value: 0.8670532608361352.

[I 2025-10-26 16:26:15,320] Trial 2 finished with value: 0.8619832915433066 and parameters: {'criterion': 'entropy', 'max_depth': 16, 'min_samples_split': 15, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8670532608361352.

[I 2025-10-26 16:26:15,710] Trial 3 finished with value: 0.8604543754886519 and parameters: {'criterion': 'gini', 'max_depth': 17, 'min_samples_split': 10, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8670532608361352.

[I 2025-10-26 16:26:16,035] Trial 4 finished with value: 0.8658457907541791 and parameters: {'criterion': 'entropy', 'max_depth': 7, 'min_samples_split': 18, 'min_samples_leaf': 7}. Best is trial 0 with value: 0.8670532608361352.

[I 2025-10-26 16:26:16,436] Trial 5 finished with value: 0.8673749981789369 and parameters: {'criterion': 'log_loss', 'max_depth': 10, 'min_samples_split': 8, 'min_samples_leaf': 1}. Best is trial 5 with value: 0.8673749981789369.

[I 2025-10-26 16:26:16,859] Trial 6 finished with value: 0.8660069184212273 and parameters: {'criterion': 'entropy', 'max_depth': 18, 'min_samples_split': 5, 'min_samples_leaf': 8}. Best is trial 5 with value: 0.8673749981789369.

[I 2025-10-26 16:26:17,328] Trial 7 finished with value: 0.8627076376197653 and parameters: {'criterion': 'entropy', 'max_depth': 14, 'min_samples_split': 19, 'min_samples_leaf': 3}. Best is trial 5 with value: 0.8673749981789369.

[I 2025-10-26 16:26:17,764] Trial 8 finished with value: 0.8630296987071262 and parameters: {'criterion': 'gini', 'max_depth': 17, 'min_samples_split': 12, 'min_samples_leaf': 6}. Best is trial 5 with value: 0.8673749981789369.

[I 2025-10-26 16:26:17,998] Trial 9 finished with value: 0.7824720163296756 and parameters: {'criterion': 'log_loss', 'max_depth': 4, 'min_samples_split': 6, 'min_samples_leaf': 10}. Best is trial 5 with value: 0.8673749981789369.

[I 2025-10-26 16:26:18,437] Trial 10 finished with value: 0.8687432074344701 and parameters: {'criterion': 'log_loss', 'max_depth': 10, 'min_samples_split': 8, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:18,855] Trial 11 finished with value: 0.8673749981789369 and parameters: {'criterion': 'log_loss', 'max_depth': 10, 'min_samples_split': 8, 'min_samples_leaf': 1}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:19,278] Trial 12 finished with value: 0.8599711867342428 and parameters: {'criterion': 'log_loss', 'max_depth': 11, 'min_samples_split': 2, 'min_samples_leaf': 1}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:19,641] Trial 13 finished with value: 0.8672947419027439 and parameters: {'criterion': 'log_loss', 'max_depth': 8, 'min_samples_split': 9, 'min_samples_leaf': 3}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:19,798] Trial 14 finished with value: 0.5961694867515633 and parameters: {'criterion': 'log_loss', 'max_depth': 2, 'min_samples_split': 5, 'min_samples_leaf': 3}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:20,234] Trial 15 finished with value: 0.8561885229316365 and parameters: {'criterion': 'log_loss', 'max_depth': 13, 'min_samples_split': 7, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:20,667] Trial 16 finished with value: 0.852567407664005 and

parameters: {'criterion': 'log_loss', 'max_depth': 20, 'min_samples_split': 3, 'min_samples_leaf': 4}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:21,030] Trial 17 finished with value: 0.8664093976570607 and parameters: {'criterion': 'log_loss', 'max_depth': 8, 'min_samples_split': 15, 'min_samples_leaf': 1}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:21,336] Trial 18 finished with value: 0.8532912033747133 and parameters: {'criterion': 'log_loss', 'max_depth': 6, 'min_samples_split': 10, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:21,753] Trial 19 finished with value: 0.8615003294100889 and parameters: {'criterion': 'log_loss', 'max_depth': 12, 'min_samples_split': 4, 'min_samples_leaf': 4}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:22,091] Trial 20 finished with value: 0.8680996032510429 and parameters: {'criterion': 'gini', 'max_depth': 10, 'min_samples_split': 8, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:22,430] Trial 21 finished with value: 0.8680996032510429 and parameters: {'criterion': 'gini', 'max_depth': 10, 'min_samples_split': 8, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8687432074344701.

[I 2025-10-26 16:26:22,749] Trial 22 finished with value: 0.8700311927882662 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 11, 'min_samples_leaf': 2}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:22,969] Trial 23 finished with value: 0.8343789850931819 and parameters: {'criterion': 'gini', 'max_depth': 5, 'min_samples_split': 12, 'min_samples_leaf': 4}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:23,287] Trial 24 finished with value: 0.8694677477576642 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 14, 'min_samples_leaf': 2}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:23,584] Trial 25 finished with value: 0.8676975448831363 and parameters: {'criterion': 'gini', 'max_depth': 8, 'min_samples_split': 15, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:23,952] Trial 26 finished with value: 0.8675361905948968 and parameters: {'criterion': 'gini', 'max_depth': 12, 'min_samples_split': 13, 'min_samples_leaf': 2}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:24,266] Trial 27 finished with value: 0.8689850122456381 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 17, 'min_samples_leaf': 4}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:24,429] Trial 28 finished with value: 0.7259775547897198 and parameters: {'criterion': 'gini', 'max_depth': 3, 'min_samples_split': 17, 'min_samples_leaf': 4}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:24,677] Trial 29 finished with value: 0.8381618755169795 and parameters: {'criterion': 'gini', 'max_depth': 6, 'min_samples_split': 17, 'min_samples_leaf': 5}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:24,995] Trial 30 finished with value: 0.8692263314399673 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 14, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:25,315] Trial 31 finished with value: 0.8692263314399673 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 14, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8700311927882662.

[I 2025-10-26 16:26:25,591] Trial 32 finished with value: 0.8653632171144323 and

parameters: {'criterion': 'gini', 'max_depth': 7, 'min_samples_split': 14, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:25,956] Trial 33 finished with value: 0.8665704282007411 and parameters: {'criterion': 'gini', 'max_depth': 12, 'min_samples_split': 13, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:26,279] Trial 34 finished with value: 0.8699506451419701 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 11, 'min_samples_leaf': 1}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:26,556] Trial 35 finished with value: 0.8654437000118167 and parameters: {'criterion': 'gini', 'max_depth': 7, 'min_samples_split': 11, 'min_samples_leaf': 1}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:26,954] Trial 36 finished with value: 0.8595690312429687 and parameters: {'criterion': 'gini', 'max_depth': 14, 'min_samples_split': 11, 'min_samples_leaf': 1}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:27,268] Trial 37 finished with value: 0.8693068467118075 and parameters: {'criterion': 'gini', 'max_depth': 9, 'min_samples_split': 16, 'min_samples_leaf': 2}. Best is trial 22 with value: 0.8700311927882662.
[I 2025-10-26 16:26:27,670] Trial 38 finished with value: 0.8717206537697626 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 16, 'min_samples_leaf': 8}. Best is trial 38 with value: 0.8717206537697626.
[I 2025-10-26 16:26:28,072] Trial 39 finished with value: 0.8726867075340214 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:28,481] Trial 40 finished with value: 0.8684210815981974 and parameters: {'criterion': 'entropy', 'max_depth': 15, 'min_samples_split': 19, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:28,878] Trial 41 finished with value: 0.8726867075340214 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:29,271] Trial 42 finished with value: 0.8726867075340214 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:29,671] Trial 43 finished with value: 0.8726867075340214 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:30,068] Trial 44 finished with value: 0.8726867075340214 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 8}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:30,473] Trial 45 finished with value: 0.8699503861463228 and parameters: {'criterion': 'entropy', 'max_depth': 13, 'min_samples_split': 20, 'min_samples_leaf': 9}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:30,882] Trial 46 finished with value: 0.8689042056036944 and parameters: {'criterion': 'entropy', 'max_depth': 13, 'min_samples_split': 19, 'min_samples_leaf': 7}. Best is trial 39 with value: 0.8726867075340214.
[I 2025-10-26 16:26:31,279] Trial 47 finished with value: 0.8727671904314057 and parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 9}. Best is trial 47 with value: 0.8727671904314057.
[I 2025-10-26 16:26:31,689] Trial 48 finished with value: 0.8683408253220044 and


```

parameters: {'criterion': 'entropy', 'max_depth': 15, 'min_samples_split': 19,
'min_samples_leaf': 10}. Best is trial 47 with value: 0.8727671904314057.
[I 2025-10-26 16:26:32,101] Trial 49 finished with value: 0.8696284869312414 and
parameters: {'criterion': 'entropy', 'max_depth': 14, 'min_samples_split': 20,
'min_samples_leaf': 9}. Best is trial 47 with value: 0.8727671904314057.
Best parameters: {'criterion': 'entropy', 'max_depth': 11, 'min_samples_split':
20, 'min_samples_leaf': 9}
Best cross-validation accuracy: 0.8727671904314057
--- Decision Tree Model Evaluation (Optimized) ---
Test Accuracy: 0.8725458641776633
Test Precision (Weighted): 0.8733532739578848
Test Recall (Weighted): 0.8725458641776633
Test F1-Score (Weighted): 0.8728714678590848

```

Confusion Matrix:

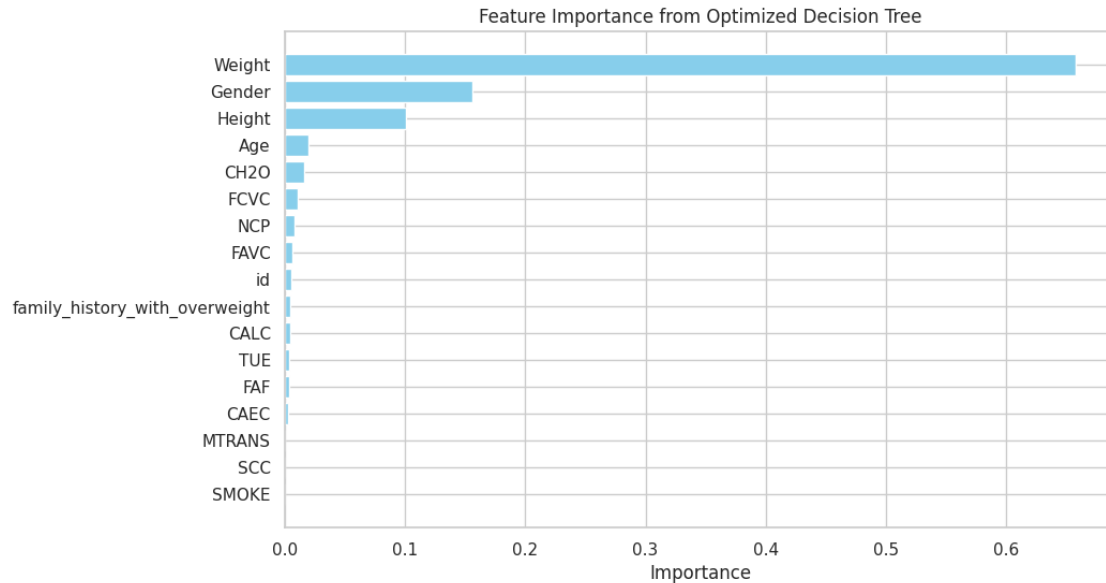
```

[[312  31   1   1   0   0   1]
 [ 27 414   2   0   0  49   4]
 [  0   1 383  16   1  16  31]
 [  0   0  24 436   0   0   4]
 [  0   0   2   1 603   0   0]
 [  3  40  10   0   0 276  55]
 [  0   6  29   3   0  38 287]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.91	0.90	0.91	346
Normal_Weight	0.84	0.83	0.84	496
Obesity_Type_I	0.85	0.85	0.85	448
Obesity_Type_II	0.95	0.94	0.95	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.73	0.72	0.72	384
Overweight_Level_II	0.75	0.79	0.77	363
accuracy			0.87	3107
macro avg	0.86	0.86	0.86	3107
weighted avg	0.87	0.87	0.87	3107



ACCURACY WITH DECISION TREE USING OPTUNA IS 0.8651432249758609

2. RANDOM FOREST

```
[57]: rf_clf = RandomForestClassifier(
    n_estimators=200,
    criterion='gini',
    max_depth=10,
    min_samples_split=5,
    min_samples_leaf=2,
    max_features='sqrt',
    random_state=42,
    n_jobs=-1
)

rf_clf.fit(X_train, y_train)

y_pred = rf_clf.predict(X_test)

test_accuracy_rf = accuracy_score(y_test, y_pred)
test_precision_rf = precision_score(y_test, y_pred, average='weighted',
    ↪zero_division=0)
test_recall_rf = recall_score(y_test, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_rf = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("--- Random Forest Model Evaluation ---")
print("Test Accuracy:", test_accuracy_rf)
```

```

print("Test Precision (Weighted):", test_precision_rf)
print("Test Recall (Weighted):", test_recall_rf)
print("Test F1-Score (Weighted):", test_f1_rf)
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

importances = rf_clf.feature_importances_
feature_names = X_train.columns
feat_importances = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'],
        color='lightgreen')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from Random Forest")
plt.show()

```

--- Random Forest Model Evaluation ---

Test Accuracy: 0.8918570968780174
 Test Precision (Weighted): 0.8934727995317859
 Test Recall (Weighted): 0.8918570968780174
 Test F1-Score (Weighted): 0.891523274112166

Confusion Matrix:

```


[[308  35   0   1   0   1   1]
 [ 19 441   1   0   0  28   7]
 [  0   1 395  13   1  10  28]
 [  0   0  12 449   0   0   3]
 [  0   0   1   1 603   1   0]
 [  2  51   8   0   0 267  56]
 [  0  12  20   1   0  22 308]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.94	0.89	0.91	346
Normal_Weight	0.82	0.89	0.85	496
Obesity_Type_I	0.90	0.88	0.89	448
Obesity_Type_II	0.97	0.97	0.97	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.81	0.70	0.75	384
Overweight_Level_II	0.76	0.85	0.80	363

accuracy			0.89	3107
macro avg	0.89	0.88	0.88	3107
weighted avg	0.89	0.89	0.89	3107



THE ACCURACY FROM RANDOM FOREST IS 0.8918570968780174

Hyperparameter Tuning Using Optuna On Random Forest

```
[58]: def objective(trial):
    n_estimators = trial.suggest_int('n_estimators', 100, 500)
    max_depth = trial.suggest_int('max_depth', 2, 20)
    min_samples_split = trial.suggest_int('min_samples_split', 2, 20)
    min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 10)
    max_features = trial.suggest_categorical('max_features', ['sqrt', 'log2',
↪None])
    criterion = trial.suggest_categorical('criterion', ['gini', 'entropy',
↪'log_loss'])

    rf_clf = RandomForestClassifier(
        n_estimators=n_estimators,
        max_depth=max_depth,
        min_samples_split=min_samples_split,
        min_samples_leaf=min_samples_leaf,
        max_features=max_features,
        criterion=criterion,
        random_state=42,
```

```

        n_jobs=-1
    )

    score = cross_val_score(rf_clf, X_train, y_train, cv=5, scoring='accuracy').
    ↪mean()
    return score

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50, show_progress_bar=True)

best_params = study.best_params
rf_clf_best = RandomForestClassifier(**best_params, random_state=42, n_jobs=-1)
rf_clf_best.fit(X_train, y_train)
y_pred = rf_clf_best.predict(X_test)

test_accuracy_rf_hp = accuracy_score(y_test, y_pred)
test_precision_rf_hp = precision_score(y_test, y_pred, average='weighted', ↪
    ↪zero_division=0)
test_recall_rf_hp = recall_score(y_test, y_pred, average='weighted', ↪
    ↪zero_division=0)
test_f1_rf_hp = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("Best parameters:", study.best_params)
print("Best cross-validation accuracy:", study.best_value)
print("--- Random Forest Model Evaluation (Optimized) ---")
print("Test Accuracy:", test_accuracy_rf_hp)
print("Test Precision (Weighted):", test_precision_rf_hp)
print("Test Recall (Weighted):", test_recall_rf_hp)
print("Test F1-Score (Weighted):", test_f1_rf_hp)
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

importances = rf_clf_best.feature_importances_
feature_names = X_train.columns
feat_importances = pd.DataFrame({'Feature': feature_names, 'Importance': ↪
    ↪importances}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'], ↪
    ↪color='lightgreen')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from Optimized Random Forest")
plt.show()

```

[I 2025-10-26 16:26:33,349] A new study created in memory with name: no-name-09421bbf-a836-42aa-aab8-94d59390f898

0%| | 0/50 [00:00<?, ?it/s]

[I 2025-10-26 16:26:46,860] Trial 0 finished with value: 0.8913566677620024 and parameters: {'n_estimators': 185, 'max_depth': 17, 'min_samples_split': 16, 'min_samples_leaf': 4, 'max_features': None, 'criterion': 'log_loss'}. Best is trial 0 with value: 0.8913566677620024.

[I 2025-10-26 16:26:54,075] Trial 1 finished with value: 0.8937714136791788 and parameters: {'n_estimators': 283, 'max_depth': 12, 'min_samples_split': 15, 'min_samples_leaf': 1, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 1 with value: 0.8937714136791788.

[I 2025-10-26 16:27:06,595] Trial 2 finished with value: 0.891195637218322 and parameters: {'n_estimators': 196, 'max_depth': 19, 'min_samples_split': 17, 'min_samples_leaf': 4, 'max_features': None, 'criterion': 'entropy'}. Best is trial 1 with value: 0.8937714136791788.

[I 2025-10-26 16:27:12,097] Trial 3 finished with value: 0.8944956626322698 and parameters: {'n_estimators': 255, 'max_depth': 14, 'min_samples_split': 7, 'min_samples_leaf': 3, 'max_features': 'sqrt', 'criterion': 'log_loss'}. Best is trial 3 with value: 0.8944956626322698.

[I 2025-10-26 16:27:15,310] Trial 4 finished with value: 0.7907624022493772 and parameters: {'n_estimators': 267, 'max_depth': 4, 'min_samples_split': 12, 'min_samples_leaf': 4, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 3 with value: 0.8944956626322698.

[I 2025-10-26 16:27:17,784] Trial 5 finished with value: 0.8894260818329123 and parameters: {'n_estimators': 125, 'max_depth': 10, 'min_samples_split': 9, 'min_samples_leaf': 6, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 3 with value: 0.8944956626322698.

[I 2025-10-26 16:27:19,391] Trial 6 finished with value: 0.7732178104831725 and parameters: {'n_estimators': 151, 'max_depth': 4, 'min_samples_split': 3, 'min_samples_leaf': 4, 'max_features': 'sqrt', 'criterion': 'gini'}. Best is trial 3 with value: 0.8944956626322698.

[I 2025-10-26 16:27:26,445] Trial 7 finished with value: 0.8959441281639962 and parameters: {'n_estimators': 320, 'max_depth': 16, 'min_samples_split': 11, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'entropy'}. Best is trial 7 with value: 0.8959441281639962.

[I 2025-10-26 16:27:31,083] Trial 8 finished with value: 0.8859652524964752 and parameters: {'n_estimators': 303, 'max_depth': 11, 'min_samples_split': 6, 'min_samples_leaf': 10, 'max_features': 'sqrt', 'criterion': 'gini'}. Best is trial 7 with value: 0.8959441281639962.

[I 2025-10-26 16:27:40,822] Trial 9 finished with value: 0.8966683447426312 and parameters: {'n_estimators': 429, 'max_depth': 15, 'min_samples_split': 4, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 9 with value: 0.8966683447426312.

[I 2025-10-26 16:27:49,131] Trial 10 finished with value: 0.8822635248336358 and parameters: {'n_estimators': 467, 'max_depth': 8, 'min_samples_split': 3, 'min_samples_leaf': 7, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 9 with value: 0.8966683447426312.

[I 2025-10-26 16:27:57,780] Trial 11 finished with value: 0.8947372084477905 and parameters: {'n_estimators': 407, 'max_depth': 16, 'min_samples_split': 20,

'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 9 with value: 0.8966683447426312.

[I 2025-10-26 16:28:05,720] Trial 12 finished with value: 0.895461619273161 and parameters: {'n_estimators': 371, 'max_depth': 20, 'min_samples_split': 12, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'entropy'}. Best is trial 9 with value: 0.8966683447426312.

[I 2025-10-26 16:28:14,398] Trial 13 finished with value: 0.8966683771170871 and parameters: {'n_estimators': 376, 'max_depth': 15, 'min_samples_split': 7, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:28:23,904] Trial 14 finished with value: 0.8891040531200073 and parameters: {'n_estimators': 493, 'max_depth': 14, 'min_samples_split': 6, 'min_samples_leaf': 10, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:28:31,259] Trial 15 finished with value: 0.8830679976884639 and parameters: {'n_estimators': 418, 'max_depth': 8, 'min_samples_split': 3, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:28:38,348] Trial 16 finished with value: 0.890954933138655 and parameters: {'n_estimators': 357, 'max_depth': 14, 'min_samples_split': 8, 'min_samples_leaf': 7, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:28:46,248] Trial 17 finished with value: 0.8959442900362756 and parameters: {'n_estimators': 435, 'max_depth': 18, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:06,627] Trial 18 finished with value: 0.8897472040610518 and parameters: {'n_estimators': 360, 'max_depth': 13, 'min_samples_split': 10, 'min_samples_leaf': 8, 'max_features': None, 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:14,924] Trial 19 finished with value: 0.8875747485718819 and parameters: {'n_estimators': 449, 'max_depth': 9, 'min_samples_split': 2, 'min_samples_leaf': 3, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:23,028] Trial 20 finished with value: 0.894254278689029 and parameters: {'n_estimators': 393, 'max_depth': 16, 'min_samples_split': 5, 'min_samples_leaf': 5, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:30,847] Trial 21 finished with value: 0.8959442900362756 and parameters: {'n_estimators': 433, 'max_depth': 18, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:39,923] Trial 22 finished with value: 0.8960246758102921 and parameters: {'n_estimators': 491, 'max_depth': 20, 'min_samples_split': 5, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 13 with value: 0.8966683771170871.

[I 2025-10-26 16:29:48,743] Trial 23 finished with value: 0.8966684742404547 and parameters: {'n_estimators': 485, 'max_depth': 20, 'min_samples_split': 8,

'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:29:56,676] Trial 24 finished with value: 0.8933689020688895 and parameters: {'n_estimators': 463, 'max_depth': 15, 'min_samples_split': 8, 'min_samples_leaf': 3, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:30:16,042] Trial 25 finished with value: 0.8900692975228687 and parameters: {'n_estimators': 342, 'max_depth': 18, 'min_samples_split': 9, 'min_samples_leaf': 1, 'max_features': None, 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:30:19,738] Trial 26 finished with value: 0.6356837404151376 and parameters: {'n_estimators': 397, 'max_depth': 2, 'min_samples_split': 13, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:30:27,854] Trial 27 finished with value: 0.8937714460536348 and parameters: {'n_estimators': 484, 'max_depth': 12, 'min_samples_split': 7, 'min_samples_leaf': 3, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:30:35,791] Trial 28 finished with value: 0.893449384966274 and parameters: {'n_estimators': 386, 'max_depth': 20, 'min_samples_split': 4, 'min_samples_leaf': 5, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:30:56,904] Trial 29 finished with value: 0.8918396622696759 and parameters: {'n_estimators': 334, 'max_depth': 17, 'min_samples_split': 10, 'min_samples_leaf': 1, 'max_features': None, 'criterion': 'log_loss'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:31:04,577] Trial 30 finished with value: 0.895059075288416 and parameters: {'n_estimators': 427, 'max_depth': 17, 'min_samples_split': 7, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:31:13,800] Trial 31 finished with value: 0.8957831299947715 and parameters: {'n_estimators': 491, 'max_depth': 20, 'min_samples_split': 4, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:31:22,585] Trial 32 finished with value: 0.8957831623692275 and parameters: {'n_estimators': 460, 'max_depth': 19, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:31:31,247] Trial 33 finished with value: 0.8965879913430704 and parameters: {'n_estimators': 475, 'max_depth': 19, 'min_samples_split': 6, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:31:55,312] Trial 34 finished with value: 0.8912763467368976 and parameters: {'n_estimators': 455, 'max_depth': 15, 'min_samples_split': 8, 'min_samples_leaf': 3, 'max_features': None, 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:04,446] Trial 35 finished with value: 0.8949785600165757 and parameters: {'n_estimators': 415, 'max_depth': 19, 'min_samples_split': 6,

'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:12,065] Trial 36 finished with value: 0.8936102860121308 and parameters: {'n_estimators': 439, 'max_depth': 13, 'min_samples_split': 9, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:17,045] Trial 37 finished with value: 0.8923226891518053 and parameters: {'n_estimators': 238, 'max_depth': 18, 'min_samples_split': 16, 'min_samples_leaf': 4, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:46,623] Trial 38 finished with value: 0.8922419472587739 and parameters: {'n_estimators': 479, 'max_depth': 15, 'min_samples_split': 7, 'min_samples_leaf': 2, 'max_features': None, 'criterion': 'log_loss'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:54,488] Trial 39 finished with value: 0.8952198468364492 and parameters: {'n_estimators': 378, 'max_depth': 17, 'min_samples_split': 13, 'min_samples_leaf': 3, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:32:58,801] Trial 40 finished with value: 0.8874944922956889 and parameters: {'n_estimators': 274, 'max_depth': 11, 'min_samples_split': 4, 'min_samples_leaf': 9, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 23 with value: 0.8966684742404547.

[I 2025-10-26 16:33:08,011] Trial 41 finished with value: 0.8967489571378392 and parameters: {'n_estimators': 500, 'max_depth': 19, 'min_samples_split': 6, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:17,268] Trial 42 finished with value: 0.8967489571378392 and parameters: {'n_estimators': 500, 'max_depth': 19, 'min_samples_split': 6, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:26,265] Trial 43 finished with value: 0.8953004268572011 and parameters: {'n_estimators': 497, 'max_depth': 16, 'min_samples_split': 7, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:34,237] Trial 44 finished with value: 0.8956222613233706 and parameters: {'n_estimators': 450, 'max_depth': 19, 'min_samples_split': 10, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:40,073] Trial 45 finished with value: 0.8532920451105668 and parameters: {'n_estimators': 471, 'max_depth': 6, 'min_samples_split': 4, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:44,730] Trial 46 finished with value: 0.8924032691725575 and parameters: {'n_estimators': 214, 'max_depth': 17, 'min_samples_split': 19, 'min_samples_leaf': 2, 'max_features': 'sqrt', 'criterion': 'entropy'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:55,867] Trial 47 finished with value: 0.8960247081847481 and parameters: {'n_estimators': 500, 'max_depth': 20, 'min_samples_split': 8,

'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:33:59,263] Trial 48 finished with value: 0.8935299002381141 and parameters: {'n_estimators': 159, 'max_depth': 13, 'min_samples_split': 6, 'min_samples_leaf': 4, 'max_features': 'log2', 'criterion': 'log_loss'}. Best is trial 41 with value: 0.8967489571378392.

[I 2025-10-26 16:34:06,649] Trial 49 finished with value: 0.895380683133394 and parameters: {'n_estimators': 407, 'max_depth': 18, 'min_samples_split': 3, 'min_samples_leaf': 2, 'max_features': 'log2', 'criterion': 'gini'}. Best is trial 41 with value: 0.8967489571378392.

Best parameters: {'n_estimators': 500, 'max_depth': 19, 'min_samples_split': 6, 'min_samples_leaf': 1, 'max_features': 'log2', 'criterion': 'gini'}

Best cross-validation accuracy: 0.8967489571378392

--- Random Forest Model Evaluation (Optimized) ---

Test Accuracy: 0.8982941744448021

Test Precision (Weighted): 0.8989410211193476

Test Recall (Weighted): 0.8982941744448021

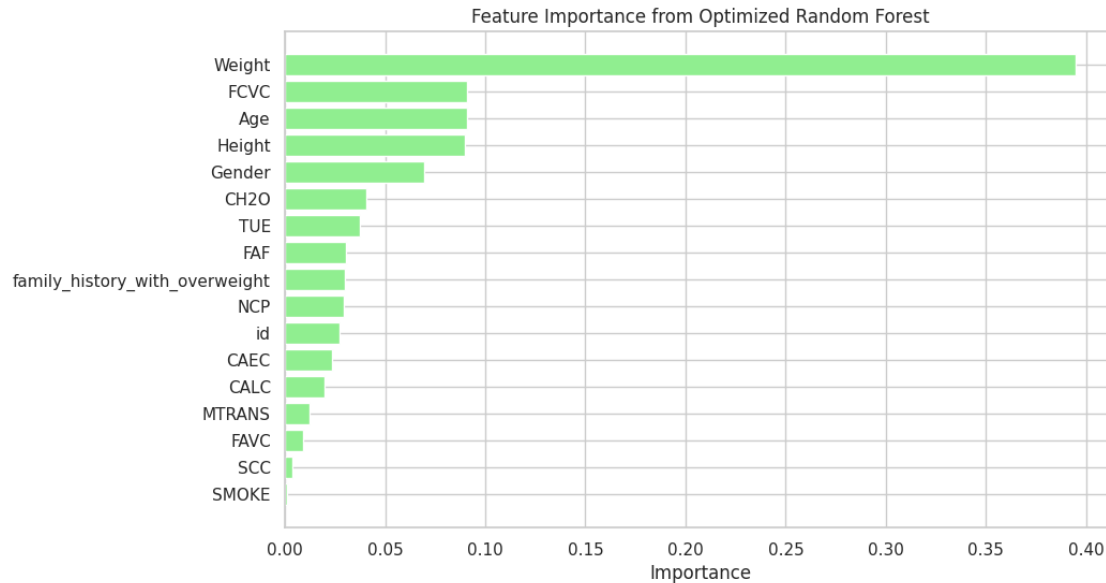
Test F1-Score (Weighted): 0.898414759101369

Confusion Matrix:

```
[[310  33   0   1   0   1   1]
 [ 20 433   2   0   0  35   6]
 [   0   1 400  13   1  10  23]
 [   0   0  12 449   0   0   3]
 [   0   0   1   1 603   1   0]
 [   2  34   8   0   0 292  48]
 [   0   9  19   2   0  29 304]]
```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.93	0.90	0.91	346
Normal_Weight	0.85	0.87	0.86	496
Obesity_Type_I	0.90	0.89	0.90	448
Obesity_Type_II	0.96	0.97	0.97	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.79	0.76	0.78	384
Overweight_Level_II	0.79	0.84	0.81	363
accuracy			0.90	3107
macro avg	0.89	0.89	0.89	3107
weighted avg	0.90	0.90	0.90	3107



THE ACCURACY FROM RANDOM FOREST USING OPTUNA IS 0.8999034438364982

3. GRADIENT BOOSTING

```
[59]: gb_clf = GradientBoostingClassifier(
        n_estimators=200,
        learning_rate=0.1,
        max_depth=3,
        min_samples_split=5,
        min_samples_leaf=2,
        subsample=0.8,
        random_state=42
    )

gb_clf.fit(X_train, y_train)

y_pred = gb_clf.predict(X_test)

test_accuracy_gb = accuracy_score(y_test, y_pred)
test_precision_gb = precision_score(y_test, y_pred, average='weighted',
    ↪zero_division=0)
test_recall_gb = recall_score(y_test, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_gb = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("--- Gradient Boosting Model Evaluation ---")
print("Test Accuracy:", test_accuracy_gb)
```

```

print("Test Precision (Weighted):", test_precision_gb)
print("Test Recall (Weighted):", test_recall_gb)
print("Test F1-Score (Weighted):", test_f1_gb)
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

importances = gb_clf.feature_importances_
feature_names = X_train.columns
feat_importances = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'],
        color='lightblue')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from Gradient Boosting")
plt.show()

```

--- Gradient Boosting Model Evaluation ---

Test Accuracy: 0.8979723205664628
 Test Precision (Weighted): 0.8982407446122279
 Test Recall (Weighted): 0.8979723205664628
 Test F1-Score (Weighted): 0.8980677084050355

Confusion Matrix:

```

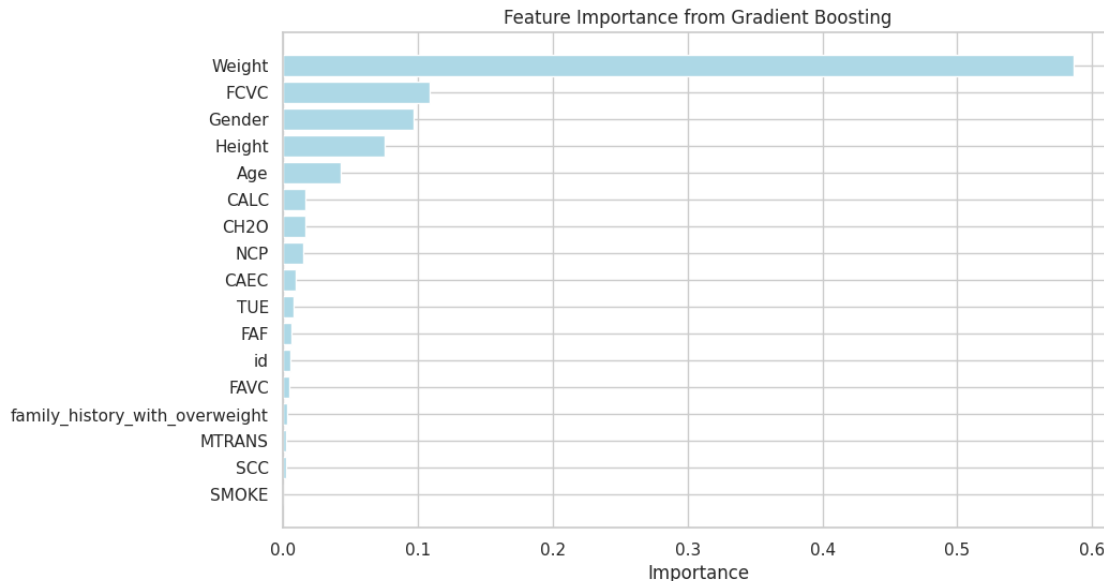
[[312  31   1   1   0   0   1]
 [ 25 432   2   0   0  34   3]
 [   1   0 397  11   1  12  26]
 [   0   0  16 448   0   0   0]
 [   0   0   2   1 603   0   0]
 [   2  34   6   0   0 299  43]
 [   0   7  21   3   0  33 299]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.92	0.90	0.91	346
Normal_Weight	0.86	0.87	0.86	496
Obesity_Type_I	0.89	0.89	0.89	448
Obesity_Type_II	0.97	0.97	0.97	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.79	0.78	0.78	384
Overweight_Level_II	0.80	0.82	0.81	363

accuracy			0.90	3107
macro avg	0.89	0.89	0.89	3107
weighted avg	0.90	0.90	0.90	3107



THE ACCURACY USING GRADIENT BOOSTING IS 0.8979723205664628

```
[60]: !pip install xgboost
```

Collecting xgboost

Downloading xgboost-3.1.1-py3-none-manylinux_2_28_x86_64.whl.metadata (2.1 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.0.2)

Collecting nvidia-nccl-cu12 (from xgboost)

Downloading nvidia_nccl_cu12-2.28.7-py3-none-manylinux_2_18_x86_64.whl.metadata (2.0 kB)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from xgboost) (1.16.2)

Downloading xgboost-3.1.1-py3-none-manylinux_2_28_x86_64.whl (115.9 MB)
115.9/115.9 MB

11.0 MB/s eta 0:00:00

Downloading nvidia_nccl_cu12-2.28.7-py3-none-manylinux_2_18_x86_64.whl (296.8 MB)

296.8/296.8 MB

1.3 MB/s eta 0:00:00

Installing collected packages: nvidia-nccl-cu12, xgboost

Successfully installed nvidia-nccl-cu12-2.28.7 xgboost-3.1.1

4. XGBOOST

```
[61]: from xgboost import XGBClassifier
```

```
[62]: y_enc = le.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(
    x, y_enc, test_size=0.2, random_state=42
)

xgb_clf = XGBClassifier(
    n_estimators=200,
    max_depth=4,
    learning_rate=0.1,
    subsample=0.8,
    colsample_bytree=0.8,
    gamma=0,
    reg_alpha=0,
    reg_lambda=1,
    random_state=42,
    eval_metric='mlogloss'
)

xgb_clf.fit(X_train, y_train)

y_pred_enc = xgb_clf.predict(X_test)

y_pred = le.inverse_transform(y_pred_enc)
y_test_orig = le.inverse_transform(y_test)

test_accuracy_xgb = accuracy_score(y_test_orig, y_pred)
test_precision_xgb = precision_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_recall_xgb = recall_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_xgb = f1_score(y_test_orig, y_pred, average='weighted', zero_division=0)

print("--- XGBoost Model Evaluation ---")
print("Test Accuracy:", test_accuracy_xgb)
print("Test Precision (Weighted):", test_precision_xgb)
print("Test Recall (Weighted):", test_recall_xgb)
print("Test F1-Score (Weighted):", test_f1_xgb)
print("\nConfusion Matrix:\n", confusion_matrix(y_test_orig, y_pred))
print("\nClassification Report:\n", classification_report(y_test_orig, y_pred))

importances = xgb_clf.feature_importances_
feature_names = X_train.columns
feat_importances = pd.DataFrame({
    'Feature': feature_names,
```

```

    'Importance': importances
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_importances['Feature'], feat_importances['Importance'],
        color='orange')
plt.gca().invert_yaxis()
plt.xlabel("Importance")
plt.title("Feature Importance from XGBoost")
plt.show()

```

```

--- XGBoost Model Evaluation ---
Test Accuracy: 0.9034438364982298
Test Precision (Weighted): 0.9039893335772461
Test Recall (Weighted): 0.9034438364982298
Test F1-Score (Weighted): 0.9036474236074407

```

Confusion Matrix:

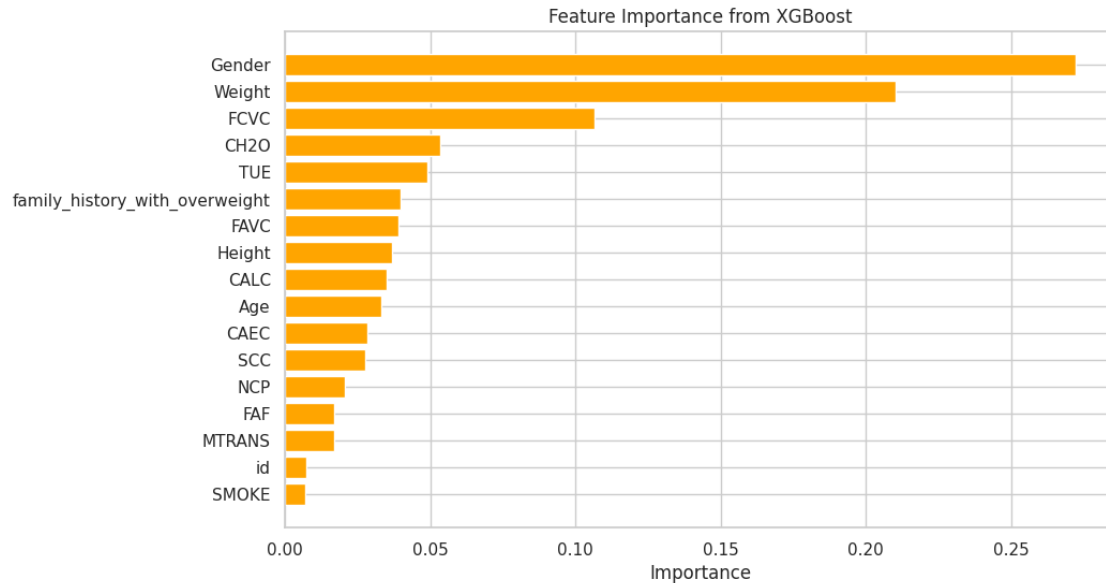
```

[[313  30   1   1   0   0   1]
 [ 19 435   2   0   0  35   5]
 [  0   0 403  10   1  12  22]
 [  0   0  11 452   0   0   1]
 [  0   0   1   1 603   0   1]
 [  2  33   4   0   0 301  44]
 [  0   7  22   1   0  33 300]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.94	0.90	0.92	346
Normal_Weight	0.86	0.88	0.87	496
Obesity_Type_I	0.91	0.90	0.90	448
Obesity_Type_II	0.97	0.97	0.97	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.79	0.78	0.79	384
Overweight_Level_II	0.80	0.83	0.81	363
accuracy			0.90	3107
macro avg	0.90	0.89	0.89	3107
weighted avg	0.90	0.90	0.90	3107



THE ACCURACY USING XGBOOST IS 0.9034438364982298

Hyperparameter Tuning Using Optuna On XGBoost

```
[63]: # =====
#   Optuna Optimization + Final XGBoost + Kaggle Submission
#   =====

import optuna
from optuna.samplers import TPESampler
from xgboost import XGBClassifier
from sklearn.model_selection import cross_val_score, StratifiedKFold, \
    train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# -----
# 1 Load Data
# -----

train_df = pd.read_csv("train.csv")
test_df = pd.read_csv("test.csv")

# Separate features and target
X = train_df.drop('WeightCategory', axis=1)
y = train_df['WeightCategory']
```



```

# Drop ID if present
if 'id' in X.columns:
    X = X.drop('id', axis=1)

test_ids = test_df['id'] if 'id' in test_df.columns else None
test_features = test_df.drop('id', axis=1, errors='ignore')

# -----
# 2 Encode categorical variables
# -----
le_target = LabelEncoder()
y_enc = le_target.fit_transform(y)

# Label encode categorical columns for both train and test
X_encoded = X.copy()
test_encoded = test_features.copy()

for col in X_encoded.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    X_encoded[col] = le.fit_transform(X_encoded[col])
    test_encoded[col] = le.transform(test_encoded[col]) # same encoder on test

# Split data for evaluation
X_train, X_valid, y_train, y_valid = train_test_split(
    X_encoded, y_enc, test_size=0.2, random_state=42, stratify=y_enc
)

print(f" Training samples: {X_train.shape[0]} | Features: {X_train.shape[1]}")

# -----
# 3 Your three best parameter sets
# -----
base_params_list = [
    {'n_estimators': 802, 'learning_rate': 0.01376727840442897, 'max_depth': 10,
    ↪ 'min_child_weight': 4, 'subsample': 0.7294665487567288, 'colsample_bytree':
    ↪ 0.5546856654158943, 'gamma': 0.16215372378274814, 'reg_alpha': 0.
    ↪ 1275266644210375, 'reg_lambda': 1.6207767848777366},

    {'n_estimators': 771, 'learning_rate': 0.013703494452479268, 'max_depth':
    ↪ 10, 'min_child_weight': 4, 'subsample': 0.7002200873954504,
    ↪ 'colsample_bytree': 0.5528037754800043, 'gamma': 0.1951947752165459,
    ↪ 'reg_alpha': 0.1397555185137831, 'reg_lambda': 1.6147146257448057},

```

```

        {'n_estimators': 809, 'learning_rate': 0.013314881450029495, 'max_depth': 9,
         'min_child_weight': 3, 'subsample': 0.7675430727257503, 'colsample_bytree': 0.5599981247033479,
         'gamma': 0.1416646576710503, 'reg_alpha': 0.12140295992038691, 'reg_lambda': 1.5505437116668446}
    ]

# Compute mean center point
mean_params = {k: np.mean([p[k] for p in base_params_list]) for k in base_params_list[0]}

# -----
# 4 Optuna objective function
# -----
def objective(trial):
    params = {
        'n_estimators': trial.suggest_int('n_estimators', 680, 820),
        'learning_rate': trial.suggest_float('learning_rate', 0.010, 0.016, log=True),
        'max_depth': trial.suggest_int('max_depth', 8, 10),
        'min_child_weight': trial.suggest_int('min_child_weight', 3, 6),
        'subsample': trial.suggest_float('subsample', 0.70, 0.80),
        'colsample_bytree': trial.suggest_float('colsample_bytree', 0.55, 0.60),
        'gamma': trial.suggest_float('gamma', 0.13, 0.20),
        'reg_alpha': trial.suggest_float('reg_alpha', 0.07, 0.14),
        'reg_lambda': trial.suggest_float('reg_lambda', 1.5, 1.7),
        'eval_metric': 'mlogloss',
        'tree_method': 'hist',
        'random_state': 42,
        'n_jobs': -1
    }

    model = XGBClassifier(**params)
    cv = StratifiedKFold(n_splits=7, shuffle=True, random_state=42)
    scores = cross_val_score(model, X_encoded, y_enc, cv=cv, scoring='accuracy', n_jobs=-1)
    return scores.mean()

# -----
# 5 Run Optuna (fewer than 100 trials)
# -----
study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=60, show_progress_bar=True)
print("Best params:", study.best_params)
print("Best accuracy:", study.best_value)

```

```

# -----
# 6 Train final model
# -----
best_params = study.best_trial.params
final_model = XGBClassifier(**best_params)
final_model.fit(X_encoded, y_enc)

# -----
# 6.1 Evaluate Final Model
# -----
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report

# Predict on validation set
y_pred_valid = final_model.predict(X_valid)

# Compute metrics
test_accuracy_xgb_hp = accuracy_score(y_valid, y_pred_valid)
test_precision_xgb_hp = precision_score(y_valid, y_pred_valid,
    average='weighted')
test_recall_xgb_hp = recall_score(y_valid, y_pred_valid, average='weighted')
test_f1_xgb_hp = f1_score(y_valid, y_pred_valid, average='weighted')
cm = confusion_matrix(y_valid, y_pred_valid)

# Display results
print("\n Model Evaluation on Validation Set")
print("=====")
print("Accuracy:", test_accuracy_xgb_hp)
print("Precision:", test_precision_xgb_hp)
print("Recall:", test_recall_xgb_hp)
print("F1 Score:", test_f1_xgb_hp)
print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_valid,
    y_pred_valid))

# -----
# 7 Make Predictions on Test Data
# -----
test_preds_enc = final_model.predict(test_encoded)
test_preds = le_target.inverse_transform(test_preds_enc)

# -----
# 8 Prepare Submission File
# -----
submission = pd.DataFrame({
    "id": test_ids,
    "WeightCategory": test_preds
})

```

```

})

submission.to_csv("submission.csv", index=False)
print("\n Submission file saved as 'submission.csv'")

# -----
# 9 Optional: Feature Importance Plot
# -----
feat_imp = pd.DataFrame({
    "Feature": X_encoded.columns,
    "Importance": final_model.feature_importances_
}).sort_values(by="Importance", ascending=False)

plt.figure(figsize=(10,6))
plt.barh(feat_imp["Feature"][:20], feat_imp["Importance"][:20],
        color="darkorange")
plt.gca().invert_yaxis()
plt.title("Top 20 Feature Importances - Final XGBoost Model")
plt.show()

```

[I 2025-10-26 16:34:58,955] A new study created in memory with name: no-name-5aa55680-6bdc-475b-98b7-52225fdc09eb

Training samples: 12426 | Features: 16

0%| | 0/60 [00:00<?, ?it/s]

[I 2025-10-26 16:35:18,549] Trial 0 finished with value: 0.9074229060709458 and parameters: {'n_estimators': 767, 'learning_rate': 0.011630836979826885, 'max_depth': 10, 'min_child_weight': 6, 'subsample': 0.7802913074065987, 'colsample_bytree': 0.5581257000430231, 'gamma': 0.15183002639795884, 'reg_alpha': 0.13868042954928156, 'reg_lambda': 1.6163301843232656}. Best is trial 0 with value: 0.9074229060709458.

[I 2025-10-26 16:35:35,052] Trial 1 finished with value: 0.9074229060709457 and parameters: {'n_estimators': 795, 'learning_rate': 0.013231395935522297, 'max_depth': 9, 'min_child_weight': 6, 'subsample': 0.7121614722908453, 'colsample_bytree': 0.5755879996239438, 'gamma': 0.1502205983473868, 'reg_alpha': 0.0850250059080261, 'reg_lambda': 1.613935063799916}. Best is trial 0 with value: 0.9074229060709458.

[I 2025-10-26 16:35:52,593] Trial 2 finished with value: 0.9080666967102299 and parameters: {'n_estimators': 768, 'learning_rate': 0.01477392524322971, 'max_depth': 10, 'min_child_weight': 4, 'subsample': 0.7074642147494602, 'colsample_bytree': 0.5571800655486081, 'gamma': 0.18105857858206562, 'reg_alpha': 0.13247485721429145, 'reg_lambda': 1.6920336068973931}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:36:08,989] Trial 3 finished with value: 0.9074229060709458 and parameters: {'n_estimators': 706, 'learning_rate': 0.010941544935194423, 'max_depth': 10, 'min_child_weight': 4, 'subsample': 0.718175497563846, 'colsample_bytree': 0.5559928939284153, 'gamma': 0.1917246458534284,

'reg_alpha': 0.10025139126185875, 'reg_lambda': 1.6636251539060338}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:36:24,539] Trial 4 finished with value: 0.9066503573038049 and parameters: {'n_estimators': 702, 'learning_rate': 0.015390765182408157, 'max_depth': 10, 'min_child_weight': 5, 'subsample': 0.7656988075043349, 'colsample_bytree': 0.5853334714001831, 'gamma': 0.175799191457729, 'reg_alpha': 0.13901504918596808, 'reg_lambda': 1.6915434754057168}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:36:43,388] Trial 5 finished with value: 0.9072297688791605 and parameters: {'n_estimators': 800, 'learning_rate': 0.015039492050992073, 'max_depth': 10, 'min_child_weight': 3, 'subsample': 0.7264218378566091, 'colsample_bytree': 0.5679230030822933, 'gamma': 0.17123185504873917, 'reg_alpha': 0.12761671247996187, 'reg_lambda': 1.5953633316988693}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:37:00,300] Trial 6 finished with value: 0.9077448013905878 and parameters: {'n_estimators': 820, 'learning_rate': 0.010985927747306435, 'max_depth': 8, 'min_child_weight': 3, 'subsample': 0.7169722254983124, 'colsample_bytree': 0.5565111770993991, 'gamma': 0.18312221486301894, 'reg_alpha': 0.10722523950021116, 'reg_lambda': 1.554281738763959}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:37:15,733] Trial 7 finished with value: 0.9073585270070171 and parameters: {'n_estimators': 726, 'learning_rate': 0.010560700241351742, 'max_depth': 8, 'min_child_weight': 3, 'subsample': 0.7511627853364868, 'colsample_bytree': 0.5931182564312805, 'gamma': 0.1713135919769499, 'reg_alpha': 0.1125422418726002, 'reg_lambda': 1.6412260312263287}. Best is trial 2 with value: 0.9080666967102299.

[I 2025-10-26 16:37:34,272] Trial 8 finished with value: 0.9081310757741583 and parameters: {'n_estimators': 789, 'learning_rate': 0.0137885644388524, 'max_depth': 10, 'min_child_weight': 3, 'subsample': 0.7383481993859331, 'colsample_bytree': 0.5588749158964819, 'gamma': 0.16596979347090798, 'reg_alpha': 0.079093945187112, 'reg_lambda': 1.6659216524948643}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:37:49,500] Trial 9 finished with value: 0.9071010107513037 and parameters: {'n_estimators': 707, 'learning_rate': 0.010961879386390034, 'max_depth': 9, 'min_child_weight': 6, 'subsample': 0.7250552074803447, 'colsample_bytree': 0.581871959328018, 'gamma': 0.13504834191013523, 'reg_alpha': 0.08852638143497504, 'reg_lambda': 1.6756209138827185}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:38:06,001] Trial 10 finished with value: 0.9072297688791605 and parameters: {'n_estimators': 744, 'learning_rate': 0.013135953991326246, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.7986786918572174, 'colsample_bytree': 0.5702593039142229, 'gamma': 0.15752652851131055, 'reg_alpha': 0.07108673090219104, 'reg_lambda': 1.5074371916627796}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:38:22,927] Trial 11 finished with value: 0.9080023176463013 and parameters: {'n_estimators': 764, 'learning_rate': 0.014212444922629104, 'max_depth': 10, 'min_child_weight': 4, 'subsample': 0.700938661408883, 'colsample_bytree': 0.5514307141832163, 'gamma': 0.19585221218596743,

'reg_alpha': 0.1228234540829291, 'reg_lambda': 1.6950996401026215}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:38:39,494] Trial 12 finished with value: 0.9078091804545163 and parameters: {'n_estimators': 782, 'learning_rate': 0.01457485850283542, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7381763609130155, 'colsample_bytree': 0.5636674925978187, 'gamma': 0.18437340228400206, 'reg_alpha': 0.07155290334453497, 'reg_lambda': 1.6503358528268959}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:38:56,813] Trial 13 finished with value: 0.907487285134874 and parameters: {'n_estimators': 742, 'learning_rate': 0.01599308018824757, 'max_depth': 10, 'min_child_weight': 3, 'subsample': 0.7525930973110482, 'colsample_bytree': 0.5617780611155483, 'gamma': 0.1622316389821933, 'reg_alpha': 0.09346052038323627, 'reg_lambda': 1.576258916341387}. Best is trial 8 with value: 0.9081310757741583.

[I 2025-10-26 16:39:13,446] Trial 14 finished with value: 0.9085817292216571 and parameters: {'n_estimators': 772, 'learning_rate': 0.013777621813274443, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.700758741960762, 'colsample_bytree': 0.5506331704907971, 'gamma': 0.18128173192348174, 'reg_alpha': 0.1180197235430194, 'reg_lambda': 1.6322892137130458}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:39:30,809] Trial 15 finished with value: 0.9083242129659436 and parameters: {'n_estimators': 820, 'learning_rate': 0.01224945331017878, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7352843509286481, 'colsample_bytree': 0.551183179328206, 'gamma': 0.14258824683050322, 'reg_alpha': 0.11719713773561674, 'reg_lambda': 1.635609819709131}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:39:48,355] Trial 16 finished with value: 0.9083242129659436 and parameters: {'n_estimators': 819, 'learning_rate': 0.012135547341759746, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7389653529966709, 'colsample_bytree': 0.5508418890459085, 'gamma': 0.1312807636876619, 'reg_alpha': 0.11812487858987344, 'reg_lambda': 1.6190130260525823}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:40:02,435] Trial 17 finished with value: 0.9070366316873751 and parameters: {'n_estimators': 685, 'learning_rate': 0.01258423537284347, 'max_depth': 8, 'min_child_weight': 5, 'subsample': 0.7699884942216029, 'colsample_bytree': 0.5502452960168863, 'gamma': 0.14568280391598226, 'reg_alpha': 0.11471749868208526, 'reg_lambda': 1.5603528431291487}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:40:19,908] Trial 18 finished with value: 0.9074229060709458 and parameters: {'n_estimators': 810, 'learning_rate': 0.011897412455849415, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7011594936820111, 'colsample_bytree': 0.5966727104948413, 'gamma': 0.13941297159912241, 'reg_alpha': 0.10353501508216349, 'reg_lambda': 1.6360667620620788}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:40:36,832] Trial 19 finished with value: 0.907165389815232 and parameters: {'n_estimators': 782, 'learning_rate': 0.013402893844101394, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.7300901972211253, 'colsample_bytree': 0.566869204875136, 'gamma': 0.19997278575643865,

'reg_alpha': 0.12241341992441451, 'reg_lambda': 1.5236728604516985}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:40:51,526] Trial 20 finished with value: 0.9072297688791605 and parameters: {'n_estimators': 731, 'learning_rate': 0.012566859813634252, 'max_depth': 8, 'min_child_weight': 5, 'subsample': 0.7607855872452357, 'colsample_bytree': 0.5745251498746516, 'gamma': 0.14364714428876812, 'reg_alpha': 0.11016289717226099, 'reg_lambda': 1.5888407245401743}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:41:08,822] Trial 21 finished with value: 0.9080023176463013 and parameters: {'n_estimators': 808, 'learning_rate': 0.011848588195725377, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7395463090325409, 'colsample_bytree': 0.5510414434663317, 'gamma': 0.1309236566974316, 'reg_alpha': 0.12062875579473815, 'reg_lambda': 1.6188302729421482}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:41:26,223] Trial 22 finished with value: 0.9081954548380866 and parameters: {'n_estimators': 819, 'learning_rate': 0.012236941894999135, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7330594642727395, 'colsample_bytree': 0.5500509710253221, 'gamma': 0.1305308608681446, 'reg_alpha': 0.11772531382010322, 'reg_lambda': 1.6311418844682866}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:41:44,023] Trial 23 finished with value: 0.9077448013905878 and parameters: {'n_estimators': 806, 'learning_rate': 0.011474265199106054, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.7433642117184942, 'colsample_bytree': 0.5542781472155464, 'gamma': 0.13829436522333158, 'reg_alpha': 0.12907725537310033, 'reg_lambda': 1.6063230138362083}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:42:00,733] Trial 24 finished with value: 0.9072297688791605 and parameters: {'n_estimators': 777, 'learning_rate': 0.010010667137984385, 'max_depth': 9, 'min_child_weight': 6, 'subsample': 0.7821106523029575, 'colsample_bytree': 0.5626032741809803, 'gamma': 0.15718044991030514, 'reg_alpha': 0.09718558815639165, 'reg_lambda': 1.652371047651585}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:42:16,839] Trial 25 finished with value: 0.9080666967102299 and parameters: {'n_estimators': 757, 'learning_rate': 0.012832721397970771, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7467943790470368, 'colsample_bytree': 0.554185301992586, 'gamma': 0.14391224539165628, 'reg_alpha': 0.1078306553969138, 'reg_lambda': 1.624392419537255}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:42:32,447] Trial 26 finished with value: 0.9080666967102298 and parameters: {'n_estimators': 793, 'learning_rate': 0.013891756897909354, 'max_depth': 8, 'min_child_weight': 4, 'subsample': 0.7226489862631477, 'colsample_bytree': 0.5604620676968505, 'gamma': 0.18870148381370716, 'reg_alpha': 0.11571888256750756, 'reg_lambda': 1.5800144179102864}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:42:49,939] Trial 27 finished with value: 0.9078735595184446 and parameters: {'n_estimators': 820, 'learning_rate': 0.012270999491364594, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7576847739367353, 'colsample_bytree': 0.5541598965752882, 'gamma': 0.13520338355589495,

'reg_alpha': 0.12869047250929894, 'reg_lambda': 1.600734779944352}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:43:05,596] Trial 28 finished with value: 0.9074229060709458 and parameters: {'n_estimators': 804, 'learning_rate': 0.013671415414302918, 'max_depth': 8, 'min_child_weight': 6, 'subsample': 0.7147310005139375, 'colsample_bytree': 0.566042062466504, 'gamma': 0.15113645428418648, 'reg_alpha': 0.12387874394435817, 'reg_lambda': 1.6461965796039095}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:43:22,952] Trial 29 finished with value: 0.9080023176463013 and parameters: {'n_estimators': 773, 'learning_rate': 0.011378105767047331, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.7768652870853008, 'colsample_bytree': 0.5728729809645302, 'gamma': 0.1771865238204776, 'reg_alpha': 0.11741722360965041, 'reg_lambda': 1.6240613084633746}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:43:39,087] Trial 30 finished with value: 0.907165389815232 and parameters: {'n_estimators': 757, 'learning_rate': 0.012945864896470614, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7315911216586517, 'colsample_bytree': 0.5812238318411875, 'gamma': 0.15757743529760676, 'reg_alpha': 0.13420269179463812, 'reg_lambda': 1.6767001618630644}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:43:58,301] Trial 31 finished with value: 0.9080666967102299 and parameters: {'n_estimators': 814, 'learning_rate': 0.012135953386674564, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7333373509774926, 'colsample_bytree': 0.5500386094082713, 'gamma': 0.1308616651361332, 'reg_alpha': 0.11889895810119998, 'reg_lambda': 1.6295886357735738}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:44:14,735] Trial 32 finished with value: 0.9074229060709458 and parameters: {'n_estimators': 796, 'learning_rate': 0.012198204785471824, 'max_depth': 9, 'min_child_weight': 6, 'subsample': 0.7067338798281085, 'colsample_bytree': 0.552803526572798, 'gamma': 0.1350427892740649, 'reg_alpha': 0.11319287820208287, 'reg_lambda': 1.6098297843565674}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:44:32,126] Trial 33 finished with value: 0.907487285134874 and parameters: {'n_estimators': 819, 'learning_rate': 0.012496981199118434, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7461059190897739, 'colsample_bytree': 0.5583480445623331, 'gamma': 0.13886129769808922, 'reg_alpha': 0.10400196159252442, 'reg_lambda': 1.6133068096704366}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:44:48,597] Trial 34 finished with value: 0.9076160432627309 and parameters: {'n_estimators': 791, 'learning_rate': 0.01184222754107582, 'max_depth': 9, 'min_child_weight': 6, 'subsample': 0.7112113716461834, 'colsample_bytree': 0.5564858641726516, 'gamma': 0.1478116947628444, 'reg_alpha': 0.13470431676178277, 'reg_lambda': 1.6335185791107714}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:45:06,062] Trial 35 finished with value: 0.9078735595184446 and parameters: {'n_estimators': 811, 'learning_rate': 0.013397555420207253, 'max_depth': 9, 'min_child_weight': 4, 'subsample': 0.7238283523559882, 'colsample_bytree': 0.5543287365672405, 'gamma': 0.13016835867751153,

'reg_alpha': 0.1260816570763907, 'reg_lambda': 1.6644693237433377}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:45:24,002] Trial 36 finished with value: 0.907487285134874 and parameters: {'n_estimators': 802, 'learning_rate': 0.011293129463408895, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7341811056964159, 'colsample_bytree': 0.557728865582524, 'gamma': 0.14115094039847892, 'reg_alpha': 0.11851857656098498, 'reg_lambda': 1.6576176030955585}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:45:43,271] Trial 37 finished with value: 0.9078735595184445 and parameters: {'n_estimators': 788, 'learning_rate': 0.013136455384422518, 'max_depth': 10, 'min_child_weight': 4, 'subsample': 0.757138860562867, 'colsample_bytree': 0.552812866397269, 'gamma': 0.13401250777986673, 'reg_alpha': 0.11062596333281233, 'reg_lambda': 1.641810553254157}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:46:00,226] Trial 38 finished with value: 0.9077448013905878 and parameters: {'n_estimators': 798, 'learning_rate': 0.01428277711522705, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7067971961600117, 'colsample_bytree': 0.5638854814147001, 'gamma': 0.15374634220455852, 'reg_alpha': 0.10007696069846139, 'reg_lambda': 1.5933736014560633}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:46:17,375] Trial 39 finished with value: 0.9074229060709457 and parameters: {'n_estimators': 814, 'learning_rate': 0.011688432243476226, 'max_depth': 9, 'min_child_weight': 5, 'subsample': 0.7203961456554931, 'colsample_bytree': 0.558970195935141, 'gamma': 0.16702648021126107, 'reg_alpha': 0.12495688424039242, 'reg_lambda': 1.678594748778006}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:46:34,586] Trial 40 finished with value: 0.907294147943089 and parameters: {'n_estimators': 718, 'learning_rate': 0.0123075920080818, 'max_depth': 10, 'min_child_weight': 4, 'subsample': 0.7418550926769639, 'colsample_bytree': 0.5896565948137471, 'gamma': 0.17925358679421355, 'reg_alpha': 0.13208596081044585, 'reg_lambda': 1.6195343693589481}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:46:52,958] Trial 41 finished with value: 0.9078735595184446 and parameters: {'n_estimators': 786, 'learning_rate': 0.013896926101093552, 'max_depth': 10, 'min_child_weight': 3, 'subsample': 0.7386147848429295, 'colsample_bytree': 0.5500322579832365, 'gamma': 0.16908099345007013, 'reg_alpha': 0.08078347024991563, 'reg_lambda': 1.6648680173658428}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:47:11,179] Trial 42 finished with value: 0.9080666967102299 and parameters: {'n_estimators': 766, 'learning_rate': 0.012895120707151323, 'max_depth': 10, 'min_child_weight': 3, 'subsample': 0.7283471737642192, 'colsample_bytree': 0.556075217289419, 'gamma': 0.16371931843320234, 'reg_alpha': 0.08066323274860213, 'reg_lambda': 1.6337599520286907}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:47:28,769] Trial 43 finished with value: 0.9085817292216571 and parameters: {'n_estimators': 820, 'learning_rate': 0.014859722339390204, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7359775951758664, 'colsample_bytree': 0.5599166707299679, 'gamma': 0.18747857936605333,

'reg_alpha': 0.10659966418097244, 'reg_lambda': 1.6865147766424577}. Best is trial 14 with value: 0.9085817292216571.

[I 2025-10-26 16:47:46,363] Trial 44 finished with value: 0.9087748664134423 and parameters: {'n_estimators': 816, 'learning_rate': 0.014710138612677975, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7180557357572586, 'colsample_bytree': 0.5522858830234695, 'gamma': 0.18769981291348767, 'reg_alpha': 0.10935646298325649, 'reg_lambda': 1.6041904508339415}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:48:03,679] Trial 45 finished with value: 0.9081954548380866 and parameters: {'n_estimators': 803, 'learning_rate': 0.015282154577699185, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7113709169962045, 'colsample_bytree': 0.5604413393759144, 'gamma': 0.18814302771919605, 'reg_alpha': 0.10685879241752422, 'reg_lambda': 1.683609662339871}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:48:21,304] Trial 46 finished with value: 0.9081954548380865 and parameters: {'n_estimators': 814, 'learning_rate': 0.01473459386467209, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.721363503127408, 'colsample_bytree': 0.5525166550189564, 'gamma': 0.19357227665154295, 'reg_alpha': 0.11039745149211673, 'reg_lambda': 1.5710620083574405}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:48:38,558] Trial 47 finished with value: 0.9087104873495139 and parameters: {'n_estimators': 798, 'learning_rate': 0.01575735720160836, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7164719182481573, 'colsample_bytree': 0.5555593098808913, 'gamma': 0.18440382043842254, 'reg_alpha': 0.10044706818651947, 'reg_lambda': 1.6035694057678818}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:48:54,732] Trial 48 finished with value: 0.9083242129659433 and parameters: {'n_estimators': 797, 'learning_rate': 0.01583107735092103, 'max_depth': 8, 'min_child_weight': 3, 'subsample': 0.7161809692127379, 'colsample_bytree': 0.5554530091778943, 'gamma': 0.18411695642716797, 'reg_alpha': 0.09259031511983017, 'reg_lambda': 1.5865439816780615}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:49:12,907] Trial 49 finished with value: 0.9068434944955899 and parameters: {'n_estimators': 777, 'learning_rate': 0.015063921369040733, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7027114039327237, 'colsample_bytree': 0.5704011368447073, 'gamma': 0.1748046828797527, 'reg_alpha': 0.09789371712692718, 'reg_lambda': 1.5416664644446408}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:49:30,802] Trial 50 finished with value: 0.9079379385823729 and parameters: {'n_estimators': 809, 'learning_rate': 0.015611341792978498, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7106780203676939, 'colsample_bytree': 0.5603426033818818, 'gamma': 0.18850596888092314, 'reg_alpha': 0.10196046626933455, 'reg_lambda': 1.6026590015475703}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:49:48,431] Trial 51 finished with value: 0.9086461082855856 and parameters: {'n_estimators': 815, 'learning_rate': 0.014250328221111297, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.726708667290631, 'colsample_bytree': 0.553164192782055, 'gamma': 0.18118846314787646,

'reg_alpha': 0.11302216140032322, 'reg_lambda': 1.6998626703145898}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:50:06,413] Trial 52 finished with value: 0.9085817292216571 and parameters: {'n_estimators': 814, 'learning_rate': 0.014434516061140375, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7275372722743714, 'colsample_bytree': 0.5567816073768916, 'gamma': 0.18215767999854976, 'reg_alpha': 0.10597588171766244, 'reg_lambda': 1.6966008919659878}. Best is trial 44 with value: 0.9087748664134423.

[I 2025-10-26 16:50:24,112] Trial 53 finished with value: 0.9088392454773707 and parameters: {'n_estimators': 814, 'learning_rate': 0.014308476269248808, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7263200600339932, 'colsample_bytree': 0.5576317868512216, 'gamma': 0.18072420150582985, 'reg_alpha': 0.10641751805672392, 'reg_lambda': 1.6915597753551928}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:50:42,055] Trial 54 finished with value: 0.9068434944955902 and parameters: {'n_estimators': 800, 'learning_rate': 0.014936519470056867, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7171132516985788, 'colsample_bytree': 0.5645389710961846, 'gamma': 0.17335715408392338, 'reg_alpha': 0.09407036449849945, 'reg_lambda': 1.6880672211176833}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:50:59,461] Trial 55 finished with value: 0.9087748664134423 and parameters: {'n_estimators': 806, 'learning_rate': 0.015223596621572908, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.719221404564782, 'colsample_bytree': 0.5588956631557571, 'gamma': 0.1866579386744173, 'reg_alpha': 0.10898606895529972, 'reg_lambda': 1.6963435237892088}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:51:17,529] Trial 56 finished with value: 0.9081310757741583 and parameters: {'n_estimators': 807, 'learning_rate': 0.015345927457087438, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.705788849313502, 'colsample_bytree': 0.5526704914280967, 'gamma': 0.17900235246355486, 'reg_alpha': 0.11291487375106399, 'reg_lambda': 1.6996907204311178}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:51:34,597] Trial 57 finished with value: 0.9078091804545163 and parameters: {'n_estimators': 785, 'learning_rate': 0.014112625434305994, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7196943840617852, 'colsample_bytree': 0.5621392683229055, 'gamma': 0.18575490085485052, 'reg_alpha': 0.10951861109901326, 'reg_lambda': 1.6931570329587995}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:51:52,138] Trial 58 finished with value: 0.9063928410480911 and parameters: {'n_estimators': 791, 'learning_rate': 0.014574943772794833, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7137923218359338, 'colsample_bytree': 0.5999426989965128, 'gamma': 0.1911333143743522, 'reg_alpha': 0.10141616314317219, 'reg_lambda': 1.668265322836546}. Best is trial 53 with value: 0.9088392454773707.

[I 2025-10-26 16:52:10,413] Trial 59 finished with value: 0.9087748664134424 and parameters: {'n_estimators': 804, 'learning_rate': 0.015553131209879384, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7241905987127736, 'colsample_bytree': 0.5550458839234195, 'gamma': 0.1975868606950675,

'reg_alpha': 0.1148692930663168, 'reg_lambda': 1.6714265650161926}. Best is trial 53 with value: 0.9088392454773707.

Best params: {'n_estimators': 814, 'learning_rate': 0.014308476269248808, 'max_depth': 9, 'min_child_weight': 3, 'subsample': 0.7263200600339932, 'colsample_bytree': 0.5576317868512216, 'gamma': 0.18072420150582985, 'reg_alpha': 0.10641751805672392, 'reg_lambda': 1.6915597753551928}

Best accuracy: 0.9088392454773707

Model Evaluation on Validation Set

=====

Accuracy: 0.964274219504345

Precision: 0.9642615720784604

Recall: 0.964274219504345

F1 Score: 0.9642404609827702

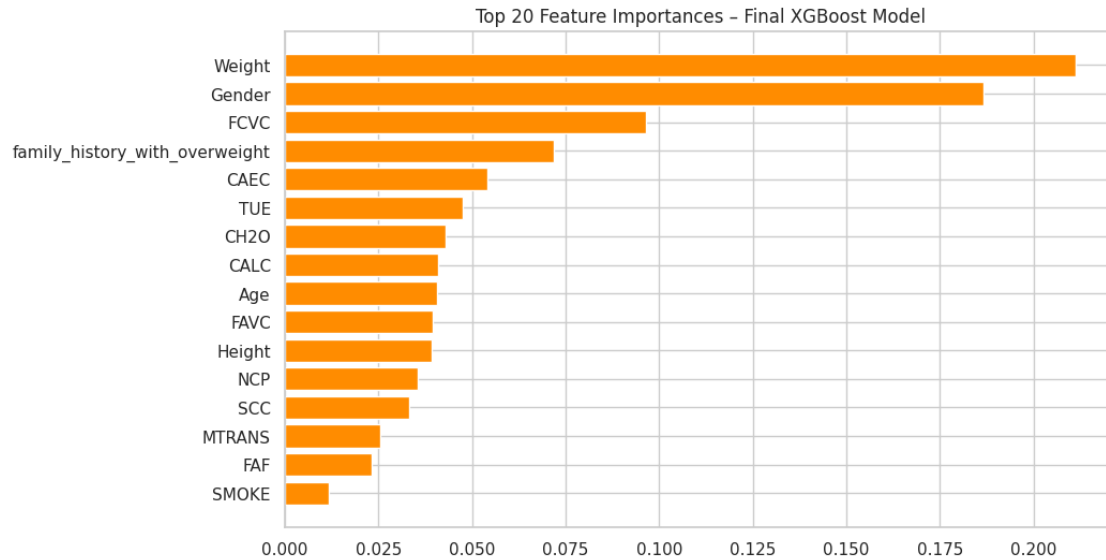
Confusion Matrix:

```
[[365   8   0   0   0   1   0]
 [  9 449   0   0   0  10   1]
 [  0   0 417   4   3   7  10]
 [  0   0   3 477   0   0   1]
 [  0   0   0   0 597   0   0]
 [  2  12   4   0   0 336  15]
 [  0   2   8   0   0  11 355]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.98	0.97	374
1	0.95	0.96	0.96	469
2	0.97	0.95	0.96	441
3	0.99	0.99	0.99	481
4	0.99	1.00	1.00	597
5	0.92	0.91	0.92	369
6	0.93	0.94	0.94	376
accuracy			0.96	3107
macro avg	0.96	0.96	0.96	3107
weighted avg	0.96	0.96	0.96	3107

Submission file saved as 'submission.csv'



XGBOOST ACCURACY USING OPTUNA IS **0.964274219504345**

5. ADABOOST

```
[71]: base_estimator = DecisionTreeClassifier(
        criterion='gini',
        max_depth=7,
        min_samples_split=5,
        min_samples_leaf=2,
        random_state=42
    )

    ada_clf = AdaBoostClassifier(
        estimator=base_estimator,
        n_estimators=100,
        learning_rate=0.5,
        random_state=42
    )

    ada_clf.fit(X_train, y_train)
    y_pred_enc = ada_clf.predict(X_test)

    y_pred = le.inverse_transform(y_pred_enc)
    y_test_orig = le.inverse_transform(y_test)

    test_accuracy_ab = accuracy_score(y_test_orig, y_pred)
    test_precision_ab = precision_score(y_test_orig, y_pred, average='weighted',
    ↪ zero_division=0)
```

```

test_recall_ab = recall_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_ab = f1_score(y_test_orig, y_pred, average='weighted', zero_division=0)

print("--- AdaBoost Model Evaluation ---")
print("Test Accuracy (AdaBoost):", test_accuracy_ab)
print("Test Precision (Weighted):", test_precision_ab)
print("Test Recall (Weighted):", test_recall_ab)
print("Test F1-Score (Weighted):", test_f1_ab)
print("\nConfusion Matrix:\n", confusion_matrix(y_test_orig, y_pred))
print("\nClassification Report:\n", classification_report(y_test_orig, y_pred))

```

```

--- AdaBoost Model Evaluation ---
Test Accuracy (AdaBoost): 0.8847763115545543
Test Precision (Weighted): 0.8860478578886273
Test Recall (Weighted): 0.8847763115545543
Test F1-Score (Weighted): 0.8852893045389144

```

Confusion Matrix:

```

[[306  37   1   1   0   0   1]
 [ 26 417   1   0   0  43   9]
 [  0   0 389  13   1  14  31]
 [  0   0  16 447   0   0   1]
 [  0   0   2   1 603   0   0]
 [  2  34   4   0   0 296  48]
 [  0   4  24   3   0  41 291]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.92	0.88	0.90	346
Normal_Weight	0.85	0.84	0.84	496
Obesity_Type_I	0.89	0.87	0.88	448
Obesity_Type_II	0.96	0.96	0.96	464
Obesity_Type_III	1.00	1.00	1.00	606
Overweight_Level_I	0.75	0.77	0.76	384
Overweight_Level_II	0.76	0.80	0.78	363
accuracy			0.88	3107
macro avg	0.88	0.87	0.88	3107
weighted avg	0.89	0.88	0.89	3107

THE ACCURACY USING ADABOOST ON DECISION TREE IS 0.8847763115545543

Hyperparameter Tuning Using Optuna On AdaBOOST

```

[ ]: def objective(trial):
    max_depth = trial.suggest_int('max_depth', 1, 10)
    min_samples_split = trial.suggest_int('min_samples_split', 2, 20)
    min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 10)
    n_estimators = trial.suggest_int('n_estimators', 50, 500)
    learning_rate = trial.suggest_float('learning_rate', 0.01, 1.0)

    base_estimator = DecisionTreeClassifier(
        criterion='gini',
        max_depth=max_depth,
        min_samples_split=min_samples_split,
        min_samples_leaf=min_samples_leaf,
        random_state=42
    )

    ada_clf = AdaBoostClassifier(
        estimator=base_estimator,
        n_estimators=n_estimators,
        learning_rate=learning_rate,
        random_state=42
    )

    score = cross_val_score(ada_clf, X_train, y_train, cv=5,
↪scoring='accuracy').mean()
    return score

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50, show_progress_bar=True)

best_params = study.best_params
base_estimator = DecisionTreeClassifier(
    criterion='gini',
    max_depth=best_params['max_depth'],
    min_samples_split=best_params['min_samples_split'],
    min_samples_leaf=best_params['min_samples_leaf'],
    random_state=42
)

ada_clf_best = AdaBoostClassifier(
    estimator=base_estimator,
    n_estimators=best_params['n_estimators'],
    learning_rate=best_params['learning_rate'],
    random_state=42
)
ada_clf_best.fit(X_train, y_train)
y_pred_enc = ada_clf_best.predict(X_test)

```

```

y_pred = le.inverse_transform(y_pred_enc)
y_test_orig = le.inverse_transform(y_test)

test_accuracy_ab_hp = accuracy_score(y_test_orig, y_pred)
test_precision_ab_hp = precision_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_recall_ab_hp = recall_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_ab_hp = f1_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)

print("Best parameters:", study.best_params)
print("Best cross-validation accuracy:", study.best_value)
print("--- AdaBoost Model Evaluation (Optimized) ---")
print("Test Accuracy (AdaBoost):", test_accuracy_ab_hp)
print("Test Precision (Weighted):", test_precision_ab_hp)
print("Test Recall (Weighted):", test_recall_ab_hp)
print("Test F1-Score (Weighted):", test_f1_ab_hp)
print("\nConfusion Matrix:\n", confusion_matrix(y_test_orig, y_pred))
print("\nClassification Report:\n", classification_report(y_test_orig, y_pred))

```

THE ACCURACY USING ADABOOST ON DECISION TREE USING OPTUNA IS 0.8937882201480528

6. KNN

```

[68]: le = LabelEncoder()
y_encoded = le.fit_transform(y)

# Now split encoded labels
X_train, X_test, y_train, y_test = train_test_split(x, y_encoded, test_size=0.
    ↪2, random_state=42)
knn_clf = KNeighborsClassifier(n_neighbors=5)

# Fit model
knn_clf.fit(X_train, y_train)

# Predict
y_pred_enc = knn_clf.predict(X_test)

# Decode predictions back to original labels
y_pred = le.inverse_transform(y_pred_enc)
y_test_orig = le.inverse_transform(y_test)

# --- Individual Classification Metrics Calculation ---
test_accuracy_knn = accuracy_score(y_test_orig, y_pred)
test_precision_knn = precision_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)

```



```

test_recall_knn = recall_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_knn = f1_score(y_test_orig, y_pred, average='weighted', zero_division=0)

# Evaluate
print("--- KNN Model Evaluation ---")
print("Test Accuracy (KNN):", test_accuracy_knn)
print("Test Precision (Weighted):", test_precision_knn)
print("Test Recall (Weighted):", test_recall_knn)
print("Test F1-Score (Weighted):", test_f1_knn)
print("\nConfusion Matrix:\n", confusion_matrix(y_test_orig, y_pred))
print("\nClassification Report:\n", classification_report(y_test_orig, y_pred))

```

--- KNN Model Evaluation ---

Test Accuracy (KNN): 0.4354682973929836

Test Precision (Weighted): 0.4292489701413928

Test Recall (Weighted): 0.4354682973929836

Test F1-Score (Weighted): 0.43105678373536

Confusion Matrix:

```

[[215 122   1   0   1   6   1]
 [148 235  11   0   0  75  27]
 [  0  12 156  38  95  68  79]
 [  0   0  34 246 181   2   1]
 [  0   0  59 240 306   0   1]
 [ 12 133  56   0   1 106  76]
 [  3  33 125   8   7  98  89]]

```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.57	0.62	0.59	346
Normal_Weight	0.44	0.47	0.46	496
Obesity_Type_I	0.35	0.35	0.35	448
Obesity_Type_II	0.46	0.53	0.49	464
Obesity_Type_III	0.52	0.50	0.51	606
Overweight_Level_I	0.30	0.28	0.29	384
Overweight_Level_II	0.32	0.25	0.28	363
accuracy			0.44	3107
macro avg	0.42	0.43	0.42	3107
weighted avg	0.43	0.44	0.43	3107

THE ACCURACY USING KNN IS 0.4354682973929836

Hyperparameter Tuning Using Optuna On KNN

```

[69]: import optuna
from sklearn.model_selection import cross_val_score

def objective(trial):
    n_neighbors = trial.suggest_int('n_neighbors', 1, 30)
    weights = trial.suggest_categorical('weights', ['uniform', 'distance'])
    metric = trial.suggest_categorical('metric', ['euclidean', 'manhattan', '
    ↪minkowski'])
    model = KNeighborsClassifier(
        n_neighbors=n_neighbors,
        weights=weights,
        metric=metric
    )
    score = cross_val_score(model, X_train, y_train, cv=5, scoring='accuracy').
    ↪mean()
    return score

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50, n_jobs=-1)

best_params = study.best_params
print("Best Parameters:", best_params)

knn_clf = KNeighborsClassifier(**best_params)
knn_clf.fit(X_train, y_train)

y_pred_enc = knn_clf.predict(X_test)
y_pred = le.inverse_transform(y_pred_enc)
y_test_orig = le.inverse_transform(y_test)

test_accuracy_knn_hp = accuracy_score(y_test_orig, y_pred)
test_precision_knn_hp = precision_score(y_test_orig, y_pred,
    ↪average='weighted', zero_division=0)
test_recall_knn_hp = recall_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)
test_f1_knn_hp = f1_score(y_test_orig, y_pred, average='weighted',
    ↪zero_division=0)

print("--- Tuned KNN Model Evaluation ---")
print("Test Accuracy (KNN):", test_accuracy_knn_hp)
print("Test Precision (Weighted):", test_precision_knn_hp)
print("Test Recall (Weighted):", test_recall_knn_hp)
print("Test F1-Score (Weighted):", test_f1_knn_hp)
print("\nConfusion Matrix:\n", confusion_matrix(y_test_orig, y_pred))
print("\nClassification Report:\n", classification_report(y_test_orig, y_pred))

```

[I 2025-10-26 16:57:35,221] A new study created in memory with name: no-

name-668063b0-867a-43d7-855a-4ea887a4455c

[I 2025-10-26 16:57:43,854] Trial 4 finished with value: 0.37035040492350724 and parameters: {'n_neighbors': 18, 'weights': 'uniform', 'metric': 'minkowski'}. Best is trial 4 with value: 0.37035040492350724.

[I 2025-10-26 16:57:44,349] Trial 21 finished with value: 0.42282166045346903 and parameters: {'n_neighbors': 5, 'weights': 'uniform', 'metric': 'euclidean'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:44,755] Trial 8 finished with value: 0.3668904496973798 and parameters: {'n_neighbors': 19, 'weights': 'uniform', 'metric': 'euclidean'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:45,648] Trial 13 finished with value: 0.3668904496973798 and parameters: {'n_neighbors': 19, 'weights': 'uniform', 'metric': 'euclidean'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:45,855] Trial 15 finished with value: 0.3520027971529903 and parameters: {'n_neighbors': 27, 'weights': 'uniform', 'metric': 'minkowski'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:46,065] Trial 22 finished with value: 0.35948683249942126 and parameters: {'n_neighbors': 21, 'weights': 'uniform', 'metric': 'minkowski'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:47,348] Trial 16 finished with value: 0.3483812286429761 and parameters: {'n_neighbors': 28, 'weights': 'uniform', 'metric': 'euclidean'}. Best is trial 21 with value: 0.42282166045346903.

[I 2025-10-26 16:57:48,256] Trial 20 finished with value: 0.48084448768232885 and parameters: {'n_neighbors': 8, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 20 with value: 0.48084448768232885.

[I 2025-10-26 16:57:48,352] Trial 11 finished with value: 0.44881562909233363 and parameters: {'n_neighbors': 19, 'weights': 'distance', 'metric': 'minkowski'}. Best is trial 20 with value: 0.48084448768232885.

[I 2025-10-26 16:57:48,847] Trial 19 finished with value: 0.4996769029300502 and parameters: {'n_neighbors': 5, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,261] Trial 12 finished with value: 0.45525299827929766 and parameters: {'n_neighbors': 14, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,460] Trial 18 finished with value: 0.4629788384368966 and parameters: {'n_neighbors': 11, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,574] Trial 10 finished with value: 0.4637039938747529 and parameters: {'n_neighbors': 9, 'weights': 'distance', 'metric': 'minkowski'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,575] Trial 6 finished with value: 0.4617719834696028 and parameters: {'n_neighbors': 13, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,655] Trial 23 finished with value: 0.4229024347209565 and parameters: {'n_neighbors': 29, 'weights': 'uniform', 'metric': 'manhattan'}. Best is trial 19 with value: 0.4996769029300502.

[I 2025-10-26 16:57:49,953] Trial 7 finished with value: 0.5087711789643735 and parameters: {'n_neighbors': 22, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 7 with value: 0.5087711789643735.

[I 2025-10-26 16:57:50,659] Trial 5 finished with value: 0.53935183101829 and parameters: {'n_neighbors': 4, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:51,547] Trial 2 finished with value: 0.5087711789643735 and parameters: {'n_neighbors': 22, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:51,758] Trial 9 finished with value: 0.5226934252336223 and parameters: {'n_neighbors': 12, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:52,267] Trial 0 finished with value: 0.5054713477971611 and parameters: {'n_neighbors': 23, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:52,553] Trial 1 finished with value: 0.5064371101913168 and parameters: {'n_neighbors': 24, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:52,758] Trial 17 finished with value: 0.5107024771314936 and parameters: {'n_neighbors': 21, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:52,958] Trial 14 finished with value: 0.5087711789643735 and parameters: {'n_neighbors': 22, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:53,457] Trial 3 finished with value: 0.5087711789643735 and parameters: {'n_neighbors': 22, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:54,848] Trial 27 finished with value: 0.35650977465759964 and parameters: {'n_neighbors': 23, 'weights': 'uniform', 'metric': 'euclidean'}. Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:55,262] Trial 24 finished with value: 0.48470957684967403 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'euclidean'}. Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:55,772] Trial 26 finished with value: 0.4791559979345097 and parameters: {'n_neighbors': 3, 'weights': 'distance', 'metric': 'minkowski'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:56,354] Trial 28 finished with value: 0.4847900597470584 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'euclidean'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:57,056] Trial 25 finished with value: 0.4229024347209565 and parameters: {'n_neighbors': 29, 'weights': 'uniform', 'metric': 'manhattan'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:57,648] Trial 30 finished with value: 0.3830663789656685 and parameters: {'n_neighbors': 13, 'weights': 'uniform', 'metric': 'minkowski'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:57:59,949] Trial 29 finished with value: 0.44776935142633756 and parameters: {'n_neighbors': 21, 'weights': 'distance', 'metric': 'minkowski'}. Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:58:01,562] Trial 31 finished with value: 0.4448718376226789 and parameters: {'n_neighbors': 26, 'weights': 'distance', 'metric': 'minkowski'}.

Best is trial 5 with value: 0.53935183101829.

[I 2025-10-26 16:58:02,258] Trial 36 finished with value: 0.5420888970184745 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 36 with value: 0.5420888970184745.

[I 2025-10-26 16:58:02,547] Trial 32 finished with value: 0.5227734872630797 and parameters: {'n_neighbors': 11, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 36 with value: 0.5420888970184745.

[I 2025-10-26 16:58:02,859] Trial 39 finished with value: 0.5422498951876991 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:02,960] Trial 38 finished with value: 0.5419270247389405 and parameters: {'n_neighbors': 3, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:03,051] Trial 37 finished with value: 0.5419270247389405 and parameters: {'n_neighbors': 3, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:04,047] Trial 33 finished with value: 0.5227734872630797 and parameters: {'n_neighbors': 11, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:04,049] Trial 34 finished with value: 0.534281667478726 and parameters: {'n_neighbors': 5, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:04,554] Trial 41 finished with value: 0.5422498951876991 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:04,560] Trial 35 finished with value: 0.5422498951876991 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:04,748] Trial 42 finished with value: 0.5420888970184745 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,254] Trial 40 finished with value: 0.5064371101913168 and parameters: {'n_neighbors': 24, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,548] Trial 46 finished with value: 0.5420888970184745 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,559] Trial 43 finished with value: 0.5422498951876991 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,650] Trial 45 finished with value: 0.5420888970184745 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,655] Trial 47 finished with value: 0.5420888970184745 and parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.

[I 2025-10-26 16:58:05,665] Trial 44 finished with value: 0.5422498951876991 and parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.

Best is trial 39 with value: 0.5422498951876991.
 [I 2025-10-26 16:58:05,758] Trial 48 finished with value: 0.5422498951876991 and
 parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric': 'manhattan'}.
 Best is trial 39 with value: 0.5422498951876991.
 [I 2025-10-26 16:58:05,798] Trial 49 finished with value: 0.5420888970184745 and
 parameters: {'n_neighbors': 1, 'weights': 'distance', 'metric': 'manhattan'}.
 Best is trial 39 with value: 0.5422498951876991.

Best Parameters: {'n_neighbors': 2, 'weights': 'distance', 'metric':
 'manhattan'}

--- Tuned KNN Model Evaluation ---

Test Accuracy (KNN): 0.5468297392983585
 Test Precision (Weighted): 0.5414058059931212
 Test Recall (Weighted): 0.5468297392983585
 Test F1-Score (Weighted): 0.5415818219119306

Confusion Matrix:

```
[[226 109   2   0   1   7   1]
 [112 234  13   1   0 104  32]
 [  0   7 166  42  75  54 104]
 [  0   0  31 291 138   0   4]
 [  0   0  16  71 519   0   0]
 [  2 108  38   0   0 132 104]
 [  2  24  96   9   7  94 131]]
```

Classification Report:

	precision	recall	f1-score	support
Insufficient_Weight	0.66	0.65	0.66	346
Normal_Weight	0.49	0.47	0.48	496
Obesity_Type_I	0.46	0.37	0.41	448
Obesity_Type_II	0.70	0.63	0.66	464
Obesity_Type_III	0.70	0.86	0.77	606
Overweight_Level_I	0.34	0.34	0.34	384
Overweight_Level_II	0.35	0.36	0.35	363
accuracy			0.55	3107
macro avg	0.53	0.53	0.52	3107
weighted avg	0.54	0.55	0.54	3107

KNN ACCURACY USING OPTUNA IS

3 CONCLUSION

```
[74]: results = [
    {"Model": "Decision Tree",
     "Accuracy": test_accuracy_dtt,
     "Precision": test_precision_dtt,
     "Recall": test_recall_dtt,
     "F1-Score": test_f1_dtt},

    {"Model": "Random Forest",
     "Accuracy": test_accuracy_rf,
     "Precision": test_precision_rf,
     "Recall": test_recall_rf,
     "F1-Score": test_f1_rf},

    {"Model": "Gradient Boosting",
     "Accuracy": test_accuracy_gb,
     "Precision": test_precision_gb,
     "Recall": test_recall_gb,
     "F1-Score": test_f1_gb},

    {"Model": "KNN",
     "Accuracy": test_accuracy_knn,
     "Precision": test_precision_knn,
     "Recall": test_recall_knn,
     "F1-Score": test_f1_knn},

    {"Model": "AdaBoost",
     "Accuracy": test_accuracy_ab,
     "Precision": test_precision_ab,
     "Recall": test_recall_ab,
     "F1-Score": test_f1_ab},

    {"Model": "XGBoost",
     "Accuracy": test_accuracy_xgb,
     "Precision": test_precision_xgb,
     "Recall": test_recall_xgb,
     "F1-Score": test_f1_xgb},
]

results_df = pd.DataFrame(results).round(4)

print(" Model Comparison DataFrame (results_df) Created:")
display(results_df)

# --- 2. Descriptive Statistics (MPA) ---
```

```

print("\n--- Descriptive Statistics of Model Performance Metrics ---")
display(results_df.set_index('Model').describe().round(4))

# --- 3. Key Insights ---
print("\n--- Key Insights from Model Metrics ---")

best_model = results_df.sort_values(by="F1-Score", ascending=False).iloc[0]
print(f" Best Overall Model (by F1-Score): **{best_model['Model']}** (F1-Score:
↳ {best_model['F1-Score']:.4f})")
print(f" Highest Accuracy Observed: {results_df['Accuracy'].max():.4f}")

# --- 4. Visualizations (MPA) ---

# Prepare data for Seaborn (melting wide format to long format)
results_melted = results_df.melt(id_vars='Model',
                                var_name='Metric',
                                value_name='Score',
                                value_vars=['Accuracy', 'Precision', 'Recall',
↳ 'F1-Score'])

# Visualization 1: Comprehensive Metric Comparison
plt.figure(figsize=(12, 7))
sns.barplot(x='Model', y='Score', hue='Metric', data=results_melted,
↳ palette='viridis')
plt.title('Comprehensive Model Metric Comparison', fontsize=16)
plt.ylabel('Score Value', fontsize=12)
plt.xlabel('Classifier Model', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(title='Metric')
plt.tight_layout()
plt.show()

# --- NEW VISUALIZATION: Accuracy Ranking Plot ---
results_accuracy_ranked = results_df.sort_values(by='Accuracy', ascending=True)

plt.figure(figsize=(10, 6))
plt.barh(results_accuracy_ranked['Model'], results_accuracy_ranked['Accuracy'],
↳ color=sns.color_palette("cividis", len(results_accuracy_ranked)))
plt.title('Models Ranked by Accuracy Score', fontsize=16)
plt.xlabel('Accuracy Score', fontsize=12)
plt.ylabel('Model', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 2: Ranked F1-Score Plot

```



```

results_ranked = results_df.sort_values(by='F1-Score', ascending=True)

plt.figure(figsize=(10, 6))
plt.barh(results_ranked['Model'], results_ranked['F1-Score'], color=sns.
    color_palette("rocket", len(results_ranked)))
plt.title('Models Ranked by F1-Score', fontsize=16)
plt.xlabel('F1-Score', fontsize=12)
plt.ylabel('Model', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```

Model Comparison DataFrame (results_df) Created:



	Model	Accuracy	Precision	Recall	F1-Score
0	Decision Tree	0.8651	0.8659	0.8651	0.8655
1	Random Forest	0.8919	0.8935	0.8919	0.8915
2	Gradient Boosting	0.8980	0.8982	0.8980	0.8981
3	KNN	0.4355	0.4292	0.4355	0.4311
4	AdaBoost	0.8848	0.8860	0.8848	0.8853
5	XGBoost	0.9034	0.9040	0.9034	0.9036


--- Descriptive Statistics of Model Performance Metrics ---

	Accuracy	Precision	Recall	F1-Score
count	6.0000	6.0000	6.0000	6.0000
mean	0.8131	0.8128	0.8131	0.8125
std	0.1855	0.1884	0.1855	0.1873
min	0.4355	0.4292	0.4355	0.4311
25%	0.8700	0.8709	0.8700	0.8704
50%	0.8884	0.8898	0.8884	0.8884
75%	0.8965	0.8970	0.8965	0.8964
max	0.9034	0.9040	0.9034	0.9036

--- Key Insights from Model Metrics ---

Best Overall Model (by F1-Score): ****XGBoost**** (F1-Score: 0.9036)
 Highest Accuracy Observed: 0.9034





```
[75]: results = [
    {"Model": "Decision Tree",
     "Accuracy": test_accuracy_dtt_hp,
     "Precision": test_precision_dtt_hp,
     "Recall": test_recall_dtt_hp,
     "F1-Score": test_f1_dtt_hp},

    {"Model": "Random Forest",
     "Accuracy": test_accuracy_rf_hp,
     "Precision": test_precision_rf_hp,
     "Recall": test_recall_rf_hp,
     "F1-Score": test_f1_rf_hp},

    {"Model": "Gradient Boosting",
     "Accuracy": test_accuracy_gb,
     "Precision": test_precision_gb,
     "Recall": test_recall_gb,
     "F1-Score": test_f1_gb},

    {"Model": "KNN",
     "Accuracy": test_accuracy_knn_hp,
     "Precision": test_precision_knn_hp,
     "Recall": test_recall_knn_hp,
     "F1-Score": test_f1_knn_hp},
```

```

        {"Model": "XGBoost",
         "Accuracy": test_accuracy_xgb_hp,
         "Precision": test_precision_xgb_hp,
         "Recall": test_recall_xgb_hp,
         "F1-Score": test_f1_xgb_hp},
    ]

results_df = pd.DataFrame(results).round(4)

print(" Model Comparison DataFrame (results_df) Created:")
display(results_df)

# --- 2. Descriptive Statistics (MPA) ---

print("\n--- Descriptive Statistics of Model Performance Metrics ---")
display(results_df.set_index('Model').describe().round(4))

# --- 3. Key Insights ---
print("\n--- Key Insights from Model Metrics ---")

best_model = results_df.sort_values(by="F1-Score", ascending=False).iloc[0]
print(f" Best Overall Model (by F1-Score): **{best_model['Model']}** (F1-Score:
    ↳ {best_model['F1-Score']:.4f})")
print(f" Highest Accuracy Observed: {results_df['Accuracy'].max():.4f}")

# --- 4. Visualizations (MPA) ---

# Prepare data for Seaborn (melting wide format to long format)
results_melted = results_df.melt(id_vars='Model',
                                var_name='Metric',
                                value_name='Score',
                                value_vars=['Accuracy', 'Precision', 'Recall',
    ↳ 'F1-Score'])

# Visualization 1: Comprehensive Metric Comparison
plt.figure(figsize=(12, 7))
sns.barplot(x='Model', y='Score', hue='Metric', data=results_melted,
    ↳ palette='viridis')
plt.title('Comprehensive Model Metric Comparison', fontsize=16)
plt.ylabel('Score Value', fontsize=12)
plt.xlabel('Classifier Model', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(title='Metric')
plt.tight_layout()
plt.show()

```

```
# --- NEW VISUALIZATION: Accuracy Ranking Plot ---
results_accuracy_ranked = results_df.sort_values(by='Accuracy', ascending=True)

plt.figure(figsize=(10, 6))
plt.barh(results_accuracy_ranked['Model'], results_accuracy_ranked['Accuracy'],
         color=sns.color_palette("cividis", len(results_accuracy_ranked)))
plt.title('Models Ranked by Accuracy Score', fontsize=16)
plt.xlabel('Accuracy Score', fontsize=12)
plt.ylabel('Model', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Visualization 2: Ranked F1-Score Plot
results_ranked = results_df.sort_values(by='F1-Score', ascending=True)

plt.figure(figsize=(10, 6))
plt.barh(results_ranked['Model'], results_ranked['F1-Score'], color=sns.
         color_palette("rocket", len(results_ranked)))
plt.title('Models Ranked by F1-Score', fontsize=16)
plt.xlabel('F1-Score', fontsize=12)
plt.ylabel('Model', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

Model Comparison DataFrame (results_df) Created:

	Model	Accuracy	Precision	Recall	F1-Score
0	Decision Tree	0.8725	0.8734	0.8725	0.8729
1	Random Forest	0.8983	0.8989	0.8983	0.8984
2	Gradient Boosting	0.8980	0.8982	0.8980	0.8981
3	KNN	0.5468	0.5414	0.5468	0.5416
4	XGBoost	0.9643	0.9643	0.9643	0.9642

--- Descriptive Statistics of Model Performance Metrics ---

	Accuracy	Precision	Recall	F1-Score
count	5.0000	5.0000	5.0000	5.0000
mean	0.8360	0.8352	0.8360	0.8350
std	0.1652	0.1677	0.1652	0.1675
min	0.5468	0.5414	0.5468	0.5416
25%	0.8725	0.8734	0.8725	0.8729
50%	0.8980	0.8982	0.8980	0.8981
75%	0.8983	0.8989	0.8983	0.8984
max	0.9643	0.9643	0.9643	0.9642

--- Key Insights from Model Metrics ---

Best Overall Model (by F1-Score): ****XGBoost**** (F1-Score: 0.9642)

Highest Accuracy Observed: 0.9643

